

Decision Tree & Random Forest V1

November 19, 2021

Replace BMI, BP, ST with mean

```
[1]: import numpy as np # Import numpy for data preprocessing
import pandas as pd # Import pandas for data frame read
import matplotlib.pyplot as plt # Import matplotlib for data visualisation
import seaborn as sns # Import seaborn for data visualisation
import plotly.express as px # Import plotly for data visualisation
from sklearn.model_selection import train_test_split # Import train_test_split
    ↳ for data split
from sklearn.tree import DecisionTreeClassifier # Import Decision Tree
    ↳ Classifier
from sklearn.ensemble import RandomForestClassifier # Import Random Forest
    ↳ Classifier
from sklearn.model_selection import train_test_split # Import train_test_split
    ↳ function
from sklearn import metrics # Import scikit-learn metrics module for accuracy
    ↳ calculation
from sklearn import tree # Import export_graphviz for visualizing Decision Trees
```

0.1 Data read

```
[2]: df = pd.read_csv("data/diabetes.csv") # Data read
```

```
[3]: df.head() # print data
```

```
[3]:
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | \ |
|---|-------------|---------|---------------|---------------|---------|------|---|
| 0 | 6 | 148 | 72 | 35 | 0 | 33.6 | |
| 1 | 1 | 85 | 66 | 29 | 0 | 26.6 | |
| 2 | 8 | 183 | 64 | 0 | 0 | 23.3 | |
| 3 | 1 | 89 | 66 | 23 | 94 | 28.1 | |
| 4 | 0 | 137 | 40 | 35 | 168 | 43.1 | |

| | DiabetesPedigreeFunction | Age | Outcome |
|---|--------------------------|-----|---------|
| 0 | 0.627 | 50 | 1 |
| 1 | 0.351 | 31 | 0 |
| 2 | 0.672 | 32 | 1 |
| 3 | 0.167 | 21 | 0 |
| 4 | 2.288 | 33 | 1 |

```
[4]: df.isna().sum() # check for null value
```

```
[4]: Pregnancies      0
      Glucose          0
      BloodPressure    0
      SkinThickness    0
      Insulin          0
      BMI              0
      DiabetesPedigreeFunction  0
      Age              0
      Outcome          0
      dtype: int64
```

```
[5]: df.describe()
```

```
[5]:
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin \ |
|-------|-------------|------------|---------------|---------------|------------|
| count | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 |
| mean | 3.845052 | 120.894531 | 69.105469 | 20.536458 | 79.799479 |
| std | 3.369578 | 31.972618 | 19.355807 | 15.952218 | 115.244002 |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 25% | 1.000000 | 99.000000 | 62.000000 | 0.000000 | 0.000000 |
| 50% | 3.000000 | 117.000000 | 72.000000 | 23.000000 | 30.500000 |
| 75% | 6.000000 | 140.250000 | 80.000000 | 32.000000 | 127.250000 |
| max | 17.000000 | 199.000000 | 122.000000 | 99.000000 | 846.000000 |

| | BMI | DiabetesPedigreeFunction | Age | Outcome |
|-------|------------|--------------------------|------------|------------|
| count | 768.000000 | 768.000000 | 768.000000 | 768.000000 |
| mean | 31.992578 | 0.471876 | 33.240885 | 0.348958 |
| std | 7.884160 | 0.331329 | 11.760232 | 0.476951 |
| min | 0.000000 | 0.078000 | 21.000000 | 0.000000 |
| 25% | 27.300000 | 0.243750 | 24.000000 | 0.000000 |
| 50% | 32.000000 | 0.372500 | 29.000000 | 0.000000 |
| 75% | 36.600000 | 0.626250 | 41.000000 | 1.000000 |
| max | 67.100000 | 2.420000 | 81.000000 | 1.000000 |

```
[6]: # replace zero bmi value with it's mean
      print("Before BMI mean : ",round(df['BMI'].mean(),1))
      df['BMI'] = df['BMI'].replace(0, df['BMI'].mean())
      print("After BMI mean : ",round(df['BMI'].mean(),1))
```

```
Before BMI mean : 32.0
```

```
After BMI mean : 32.5
```

```
[7]: # replace zero skinthickness value with it's mean
      print("Before SkinThickness mean : ",round(df['SkinThickness'].mean(),1))
      df['SkinThickness'] = df['SkinThickness'].replace(0, df['SkinThickness'].mean())
      print("After SkinThickness mean : ",round(df['SkinThickness'].mean(),1))
```

Before SkinThickness mean : 20.5

After SkinThickness mean : 26.6

```
[8]: # replace zero bloodpressure value with it's mean
print("Before BloodPressure mean : ",round(df['BloodPressure'].mean(),1))
df['BloodPressure'] = df['BloodPressure'].replace(0, df['BloodPressure'].mean())
print("After BloodPressure mean : ",round(df['BloodPressure'].mean(),1))
```

Before BloodPressure mean : 69.1

After BloodPressure mean : 72.3

```
[9]: df.describe()
```

```
[9]:
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin \ |
|-------|-------------|------------|---------------|---------------|------------|
| count | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 |
| mean | 3.845052 | 120.894531 | 72.254807 | 26.606479 | 79.799479 |
| std | 3.369578 | 31.972618 | 12.115932 | 9.631241 | 115.244002 |
| min | 0.000000 | 0.000000 | 24.000000 | 7.000000 | 0.000000 |
| 25% | 1.000000 | 99.000000 | 64.000000 | 20.536458 | 0.000000 |
| 50% | 3.000000 | 117.000000 | 72.000000 | 23.000000 | 30.500000 |
| 75% | 6.000000 | 140.250000 | 80.000000 | 32.000000 | 127.250000 |
| max | 17.000000 | 199.000000 | 122.000000 | 99.000000 | 846.000000 |

| | BMI | DiabetesPedigreeFunction | Age | Outcome |
|-------|------------|--------------------------|------------|------------|
| count | 768.000000 | 768.000000 | 768.000000 | 768.000000 |
| mean | 32.450805 | 0.471876 | 33.240885 | 0.348958 |
| std | 6.875374 | 0.331329 | 11.760232 | 0.476951 |
| min | 18.200000 | 0.078000 | 21.000000 | 0.000000 |
| 25% | 27.500000 | 0.243750 | 24.000000 | 0.000000 |
| 50% | 32.000000 | 0.372500 | 29.000000 | 0.000000 |
| 75% | 36.600000 | 0.626250 | 41.000000 | 1.000000 |
| max | 67.100000 | 2.420000 | 81.000000 | 1.000000 |

```
[10]: df.corr()
```

```
[10]:
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness \ |
|--------------------------|-------------|----------|---------------|-----------------|
| Pregnancies | 1.000000 | 0.129459 | 0.208984 | 0.013376 |
| Glucose | 0.129459 | 1.000000 | 0.218579 | 0.145378 |
| BloodPressure | 0.208984 | 0.218579 | 1.000000 | 0.134155 |
| SkinThickness | 0.013376 | 0.145378 | 0.134155 | 1.000000 |
| Insulin | -0.073535 | 0.331357 | -0.038147 | 0.286469 |
| BMI | 0.021546 | 0.218814 | 0.281231 | 0.535703 |
| DiabetesPedigreeFunction | -0.033523 | 0.137337 | 0.000371 | 0.154961 |
| Age | 0.544341 | 0.263514 | 0.326740 | 0.026423 |
| Outcome | 0.221898 | 0.466581 | 0.162986 | 0.175026 |

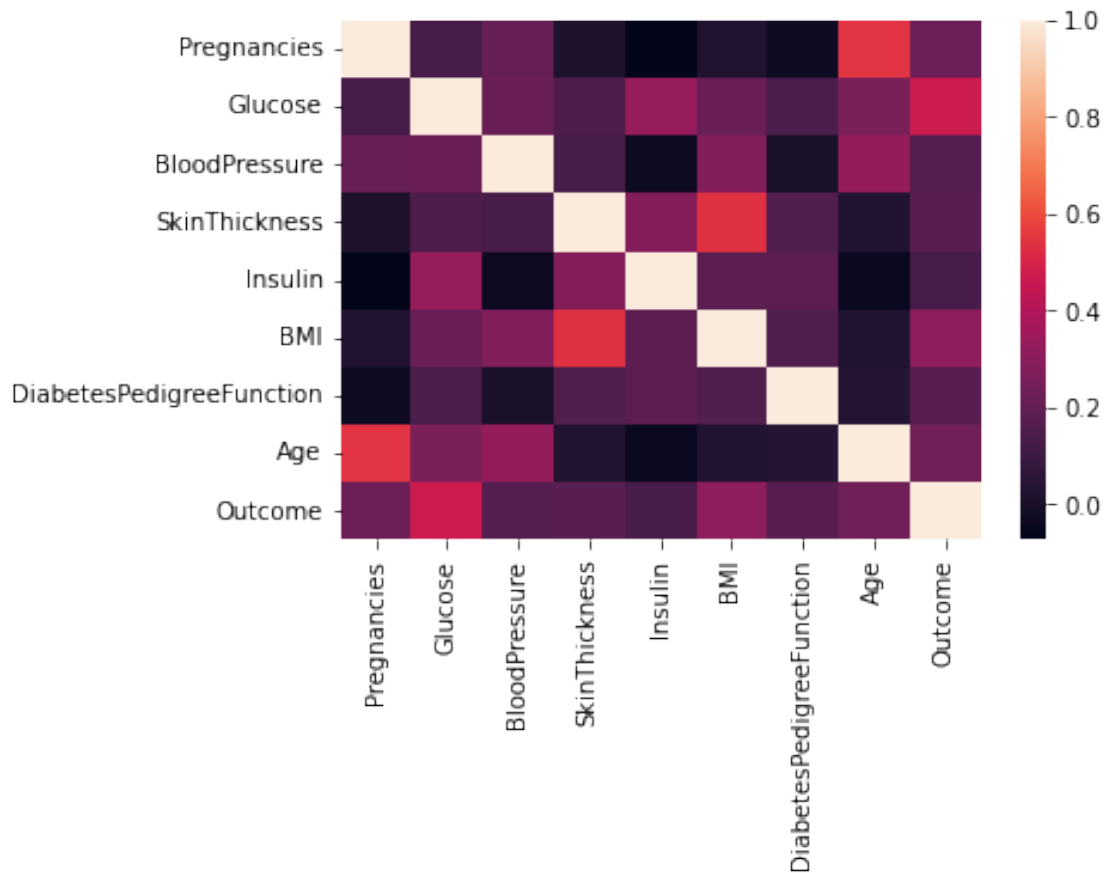
| | Insulin | BMI | DiabetesPedigreeFunction \ |
|-------------|-----------|----------|----------------------------|
| Pregnancies | -0.073535 | 0.021546 | -0.033523 |

| | | | |
|--------------------------|-----------|----------|----------|
| Glucose | 0.331357 | 0.218814 | 0.137337 |
| BloodPressure | -0.038147 | 0.281231 | 0.000371 |
| SkinThickness | 0.286469 | 0.535703 | 0.154961 |
| Insulin | 1.000000 | 0.185365 | 0.185071 |
| BMI | 0.185365 | 1.000000 | 0.153508 |
| DiabetesPedigreeFunction | 0.185071 | 0.153508 | 1.000000 |
| Age | -0.042163 | 0.025748 | 0.033561 |
| Outcome | 0.130548 | 0.312254 | 0.173844 |

| | Age | Outcome |
|--------------------------|-----------|----------|
| Pregnancies | 0.544341 | 0.221898 |
| Glucose | 0.263514 | 0.466581 |
| BloodPressure | 0.326740 | 0.162986 |
| SkinThickness | 0.026423 | 0.175026 |
| Insulin | -0.042163 | 0.130548 |
| BMI | 0.025748 | 0.312254 |
| DiabetesPedigreeFunction | 0.033561 | 0.173844 |
| Age | 1.000000 | 0.238356 |
| Outcome | 0.238356 | 1.000000 |

```
[11]: sns.heatmap(df.corr())
```

```
[11]: <AxesSubplot:>
```



1 Data split

```
[12]: X = df.iloc[:,0:-1] # All features
      Y = df.iloc[:, -1] # Target
```

```
[13]: X.head()
```

```
[13]: Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  \
0             6      148           72.0      35.000000      0  33.6
1             1       85           66.0      29.000000      0  26.6
2             8      183           64.0      20.536458      0  23.3
3             1       89           66.0      23.000000     94  28.1
4             0      137           40.0      35.000000    168  43.1

      DiabetesPedigreeFunction  Age
0                0.627      50
1                0.351      31
2                0.672      32
```

| | | |
|---|-------|----|
| 3 | 0.167 | 21 |
| 4 | 2.288 | 33 |

```
[14]: Y.head()
```

```
[14]: 0    1
      1    0
      2    1
      3    0
      4    1
      Name: Outcome, dtype: int64
```

```
[15]: # Data split
x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.2,
    random_state=1)
# x_dev, x_test, y_dev, y_test = train_test_split(x_test, y_test, test_size=0.
    5)
```

```
[16]: print("Original data size : ", X.shape, Y.shape)
      print("Train data size : ", x_train.shape, y_train.shape)
      # print("Dev data size : ", x_dev.shape, y_dev.shape)
      print("Test data size : ", x_test.shape, y_test.shape)
```

```
Original data size : (768, 8) (768,)
Train data size : (614, 8) (614,)
Test data size : (154, 8) (154,)
```

2 Decision Tree

```
[17]: accuracy = {}
```

2.0.1 criterion="gini", splitter="best"

```
[18]: # Define and build model
      clf = DecisionTreeClassifier(criterion="gini", splitter="best")
      clf = clf.fit(x_train, y_train)
      y_pred = clf.predict(x_test)
```

```
[19]: print(y_pred)
```

```
[0 0 0 1 0 0 1 0 0 0 1 0 1 1 0 0 0 1 0 1 0 0 1 0 0 1 1 1 0 1 0 0 0 1 0 1 0
 0 0 0 0 0 0 1 0 0 0 1 0 0 0 1 1 0 1 0 1 0 0 1 1 0 1 0 1 0 0 0 1 1 1 0 1 0
 0 0 1 1 0 1 1 0 0 1 0 0 1 1 1 1 0 0 0 0 1 1 0 0 1 0 0 0 1 0 0 0 0 1 0 0 0
 0 0 1 0 0 0 0 0 1 1 1 0 1 0 0 1 0 0 0 1 0 1 1 0 0 0 1 0 0 1 0 0 1 1 1 0 0
 1 0 0 1 1 0]
```

```
[20]: print(np.array(y_test))
```

```
[0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1 1 0 0 0 1 1 1 1 0 0 0 1 0 1 1 0 0 1 0 1 0
0 0 0 0 0 0 1 0 0 1 1 0 1 0 0 1 0 1 0 1 0 0 0 0 0 1 0 1 0 1 1 0 1 1 0 0 0
1 1 1 0 0 1 1 0 1 1 0 0 1 0 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 1 0 1 0 1
0 0 0 1 0 0 1 0 1 0 1 1 0 0 0 0 1 0 0 1 0 1 0 0 0 0 0 1 0 1 0 0 1 1 1 0 0
1 0 0 1 0 0]
```

```
[21]: accuracy["dt_gini_best"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:", metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.7012987012987013

```
[22]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[74 25]
 [21 34]]
```

```
[23]: print(metrics.classification_report(y_test, y_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.78 | 0.75 | 0.76 | 99 |
| 1 | 0.58 | 0.62 | 0.60 | 55 |
| accuracy | | | 0.70 | 154 |
| macro avg | 0.68 | 0.68 | 0.68 | 154 |
| weighted avg | 0.71 | 0.70 | 0.70 | 154 |

2.0.2 criterion="gini", splitter="best", max_depth=8

```
[24]: # Define and build model
clf = DecisionTreeClassifier(criterion="gini", splitter="best", max_depth=8)
clf = clf.fit(x_train, y_train)
y_pred = clf.predict(x_test)
```

```
[25]: print(y_pred)
```

```
[0 0 0 1 0 0 1 0 0 0 1 0 1 0 0 0 0 0 0 0 0 0 0 1 0 1 1 0 0 0 0 1 0 1 0
0 0 0 0 0 0 1 0 0 0 1 0 0 0 0 1 0 1 0 0 0 0 1 1 0 0 0 1 0 0 1 1 1 1 1 0 0
0 0 1 1 0 1 1 0 0 0 0 0 0 1 1 1 0 0 0 0 1 1 0 0 1 0 0 0 0 1 0 0 0 1 0 0 0
0 0 0 0 0 0 0 0 1 0 1 0 0 0 0 1 0 0 0 0 0 1 0 0 0 0 1 0 0 1 0 0 1 0 0 0 0
0 0 0 1 1 0]
```

```
[26]: print(np.array(y_test))
```

```
[0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1 1 0 0 0 1 1 1 1 0 0 0 1 0 1 1 0 0 1 0 1 0
0 0 0 0 0 0 1 0 0 1 1 0 1 0 0 1 0 1 0 1 0 0 0 0 0 1 0 1 0 1 1 0 1 1 0 0 0
1 1 1 0 0 1 1 0 1 1 0 0 1 0 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 1 0 1 0 1
0 0 0 1 0 0 1 0 1 0 1 1 0 0 0 0 1 0 0 1 0 1 0 0 0 0 0 1 0 1 0 0 1 1 1 0 0
1 0 0 1 0 0]
```

```
[27]: accuracy["dt_gini_best_8"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:", metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.6688311688311688

```
[28]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[80 19]
 [32 23]]
```

```
[29]: print(metrics.classification_report(y_test, y_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.71 | 0.81 | 0.76 | 99 |
| 1 | 0.55 | 0.42 | 0.47 | 55 |
| accuracy | | | 0.67 | 154 |
| macro avg | 0.63 | 0.61 | 0.62 | 154 |
| weighted avg | 0.65 | 0.67 | 0.66 | 154 |

2.0.3 criterion="entropy", splitter="best"

```
[30]: # Define and build model
clf = DecisionTreeClassifier(criterion="entropy", splitter="best")
clf = clf.fit(x_train, y_train)
y_pred = clf.predict(x_test)
```

```
[31]: print(y_pred)
```

```
[1 0 0 0 0 0 0 0 0 0 0 1 1 0 1 0 1 0 0 1 0 1 0 0 0 1 0 1 0 0 0 1 0 1 1 1 0
 1 0 0 0 0 0 1 0 0 1 0 0 0 0 1 1 0 1 0 0 0 0 0 1 0 0 0 0 0 0 1 1 1 1 1 0 0
 1 1 1 1 0 1 1 1 0 0 0 0 0 1 1 1 0 0 1 0 1 1 0 1 1 0 0 0 1 0 0 0 1 0 0 0 1
 0 0 1 0 0 1 0 1 1 0 1 0 0 0 0 1 0 1 0 1 0 1 1 0 0 0 0 0 0 1 0 0 1 1 1 0 0
 0 1 0 1 1 1]
```

```
[32]: print(np.array(y_test))
```

```
[0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1 1 0 0 0 1 1 1 1 0 0 0 1 0 1 1 0 0 1 0 1 0
 0 0 0 0 0 0 1 0 0 1 1 0 1 0 0 1 0 1 0 1 0 0 0 0 0 1 0 1 0 1 1 0 1 1 0 0 0
 1 1 1 0 0 1 1 0 1 1 0 0 1 0 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 1 0 1 0 1
 0 0 0 1 0 0 1 0 1 0 1 1 0 0 0 0 1 0 0 1 0 1 0 0 0 0 0 1 0 1 0 0 1 1 1 0 0
 1 0 0 1 0 0]
```

```
[33]: accuracy["dt_entropy_best"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:", metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.6688311688311688


```
[34]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[70 29]
 [22 33]]
```

```
[35]: print(metrics.classification_report(y_test, y_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.76 | 0.71 | 0.73 | 99 |
| 1 | 0.53 | 0.60 | 0.56 | 55 |
| accuracy | | | 0.67 | 154 |
| macro avg | 0.65 | 0.65 | 0.65 | 154 |
| weighted avg | 0.68 | 0.67 | 0.67 | 154 |

2.0.4 criterion="entropy", splitter="best", max_depth=8

```
[36]: # Define and build model
      clf = DecisionTreeClassifier(criterion="entropy", splitter="best", max_depth=8)
      clf = clf.fit(x_train,y_train)
      y_pred = clf.predict(x_test)
```

```
[37]: print(y_pred)
```

```
[1 0 0 0 0 0 1 0 0 0 1 1 0 1 0 0 0 1 0 0 0 0 1 1 0 0 0 1 0 0 0 1 0 1 0 0 0
 0 0 0 0 0 0 1 0 0 1 0 0 0 0 0 1 0 1 0 0 0 0 0 1 0 1 0 0 0 0 1 0 1 1 1 0 0
 1 0 1 0 0 1 1 0 0 0 0 0 1 1 1 1 0 0 0 0 1 1 0 1 1 0 0 0 1 0 0 0 0 0 0 0 1
 0 0 0 0 0 1 0 0 1 0 1 0 0 0 0 0 0 1 0 1 0 1 1 0 0 0 1 0 0 1 0 0 0 1 1 0 0
 0 1 0 1 1 1]
```

```
[38]: print(np.array(y_test))
```

```
[0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1 1 0 0 0 1 1 1 1 0 0 0 1 0 1 1 0 0 1 0 1 0
 0 0 0 0 0 0 1 0 0 1 1 0 1 0 0 1 0 1 0 1 0 0 0 0 0 1 0 1 0 1 1 0 1 1 0 0 0
 1 1 1 0 0 1 1 0 1 1 0 0 1 0 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 1 0 1 0 1
 0 0 0 1 0 0 1 0 1 0 1 1 0 0 0 0 1 0 0 1 0 1 0 0 0 0 0 1 0 1 0 0 1 1 1 0 0
 1 0 0 1 0 0]
```

```
[39]: accuracy["dt_entropy_best_8"] = metrics.accuracy_score(y_test, y_pred);
      print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.7077922077922078

```
[40]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[79 20]
 [25 30]]
```

```
[41]: print(metrics.classification_report(y_test, y_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.76 | 0.80 | 0.78 | 99 |
| 1 | 0.60 | 0.55 | 0.57 | 55 |
| accuracy | | | 0.71 | 154 |
| macro avg | 0.68 | 0.67 | 0.67 | 154 |
| weighted avg | 0.70 | 0.71 | 0.70 | 154 |

2.0.5 criterion="entropy", splitter="random"

```
[42]: # Define and build model
      clf = DecisionTreeClassifier(criterion="entropy", splitter="random")
      clf = clf.fit(x_train,y_train)
      y_pred = clf.predict(x_test)
```

```
[43]: print(y_pred)
```

```
[0 1 0 0 0 1 1 0 1 0 1 0 0 1 0 0 1 1 0 0 0 0 0 0 0 1 0 1 0 0 0 0 0 1 1 1 0
 1 0 1 0 0 0 1 0 0 1 0 0 0 0 0 0 0 1 0 0 0 1 1 0 0 1 0 0 0 0 1 1 1 1 0 1 0
 0 0 1 1 0 0 1 0 0 0 0 1 1 1 0 1 0 0 1 0 1 0 1 0 1 0 0 0 1 0 0 0 0 1 0 1 0
 0 0 0 0 0 0 0 0 1 0 1 0 0 0 1 1 0 1 1 1 0 0 0 0 0 0 1 0 0 1 0 0 0 1 0 0 0
 1 0 1 1 0 0]
```

```
[44]: print(np.array(y_test))
```

```
[0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1 1 0 0 0 1 1 1 1 0 0 0 1 0 1 1 0 0 1 0 1 0
 0 0 0 0 0 0 1 0 0 1 1 0 1 0 0 1 0 1 0 1 0 0 0 0 0 1 0 1 0 1 1 0 1 1 0 0 0
 1 1 1 0 0 1 1 0 1 1 0 0 1 0 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 1 0 1 0 1
 0 0 0 1 0 0 1 0 1 0 1 1 0 0 0 0 1 0 0 1 0 1 0 0 0 0 0 1 0 1 0 0 1 1 1 0 0
 1 0 0 1 0 0]
```

```
[45]: accuracy["dt_entropy_random"] = metrics.accuracy_score(y_test, y_pred);
      print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.6233766233766234

```
[46]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[71 28]
 [30 25]]
```

```
[47]: print(metrics.classification_report(y_test, y_pred))
```

| | precision | recall | f1-score | support |
|---|-----------|--------|----------|---------|
| 0 | 0.70 | 0.72 | 0.71 | 99 |

| | | | | | |
|--------------|------|------|------|------|-----|
| | 1 | 0.47 | 0.45 | 0.46 | 55 |
| accuracy | | | | 0.62 | 154 |
| macro avg | 0.59 | 0.59 | 0.59 | | 154 |
| weighted avg | 0.62 | 0.62 | 0.62 | | 154 |

2.0.6 criterion="entropy", splitter="random", max_depth=8

```
[48]: # Define and build model
      clf = DecisionTreeClassifier(criterion="entropy", splitter="random",
      ↪max_depth=8)
      clf = clf.fit(x_train,y_train)
      y_pred = clf.predict(x_test)
```

```
[49]: print(y_pred)
```

```
[1 0 0 1 0 0 0 0 1 0 1 0 1 1 0 1 1 0 0 0 0 0 1 0 0 1 0 0 0 1 0 0 0 1 0 1 0
 0 0 1 0 0 0 1 0 0 1 1 0 0 0 1 1 0 1 0 1 0 0 1 1 0 1 0 1 0 0 0 1 1 1 0 0 0
 1 0 1 1 0 0 1 0 0 0 0 1 1 1 0 1 1 0 0 0 1 1 1 0 1 0 0 0 1 0 0 0 1 1 0 0 1
 0 1 0 0 0 1 0 0 1 0 1 0 1 0 0 0 0 1 0 1 0 0 0 1 0 0 1 0 0 1 0 1 1 1 1 0 0
 1 1 0 1 0 0]
```

```
[50]: print(np.array(y_test))
```

```
[0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1 1 0 0 0 1 1 1 1 0 0 0 1 0 1 1 0 0 1 0 1 0
 0 0 0 0 0 0 1 0 0 1 1 0 1 0 0 1 0 1 0 1 0 0 0 0 0 1 0 1 0 1 1 0 1 1 0 0 0
 1 1 1 0 0 1 1 0 1 1 0 0 1 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 1 0 1 0 1
 0 0 0 1 0 0 1 0 1 0 1 1 0 0 0 0 1 0 0 1 0 1 0 0 0 0 0 1 0 1 0 0 1 1 1 0 0
 1 0 0 1 0 0]
```

```
[51]: accuracy["dt_entropy_random_8"] = metrics.accuracy_score(y_test, y_pred);
      print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.7077922077922078

```
[52]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[73 26]
 [19 36]]
```

```
[53]: print(metrics.classification_report(y_test, y_pred))
```

| | precision | recall | f1-score | support |
|-----------|-----------|--------|----------|---------|
| 0 | 0.79 | 0.74 | 0.76 | 99 |
| 1 | 0.58 | 0.65 | 0.62 | 55 |
| accuracy | | | 0.71 | 154 |
| macro avg | 0.69 | 0.70 | 0.69 | 154 |

| | | | | |
|--------------|------|------|------|-----|
| weighted avg | 0.72 | 0.71 | 0.71 | 154 |
|--------------|------|------|------|-----|

2.0.7 criterion="entropy", splitter="best", max_depth=3

```
[54]: # Define and build model
      clf = DecisionTreeClassifier(criterion="entropy", splitter="best", max_depth=3)
      clf = clf.fit(x_train,y_train)
      y_pred = clf.predict(x_test)
```

```
[55]: print(y_pred)
```

```
[0 0 0 0 0 0 0 0 0 0 0 1 0 1 1 0 1 0 0 0 0 0 0 1 0 0 1 0 1 0 1 0 0 0 1 0 1 0
 0 0 0 0 0 0 1 0 0 1 1 0 0 0 0 1 0 0 0 1 0 0 0 1 0 1 0 0 0 0 0 1 1 1 1 0 0
 1 0 1 0 0 1 1 0 0 1 0 1 1 0 1 0 0 0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 1 1 0 0 1
 0 1 0 0 0 0 0 0 1 0 1 0 0 0 0 0 0 0 0 0 1 0 1 1 0 0 0 1 0 0 1 0 0 1 1 0 0 0
 0 0 0 1 1 0]
```

```
[56]: print(np.array(y_test))
```

```
[0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1 1 0 0 0 1 1 1 1 0 0 0 1 0 1 1 0 0 1 0 1 0
 0 0 0 0 0 0 1 0 0 1 1 0 1 0 0 1 0 1 0 1 0 0 0 0 0 1 0 1 0 1 1 0 1 1 0 0 0
 1 1 1 0 0 1 1 0 1 1 0 0 1 0 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 1 0 1 0 1
 0 0 0 1 0 0 1 0 1 0 1 1 0 0 0 0 1 0 0 1 0 1 0 0 0 0 0 1 0 1 0 0 1 1 1 0 0
 1 0 0 1 0 0]
```

```
[57]: accuracy["dt_entropy_best_3"] = metrics.accuracy_score(y_test, y_pred);
      print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.7922077922077922

```
[58]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[87 12]
 [20 35]]
```

```
[59]: print(metrics.classification_report(y_test, y_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.81 | 0.88 | 0.84 | 99 |
| 1 | 0.74 | 0.64 | 0.69 | 55 |
| accuracy | | | 0.79 | 154 |
| macro avg | 0.78 | 0.76 | 0.77 | 154 |
| weighted avg | 0.79 | 0.79 | 0.79 | 154 |

```
[60]: feature_imp = pd.Series(clf.feature_importances_,index=X.columns).
      ↪sort_values(ascending=False)
```

```

print(feature_imp)
# Creating a bar plot
sns.barplot(x=feature_imp, y=feature_imp.index)
# Add labels to your graph
plt.xlabel('Feature Importance Score')
plt.ylabel('Features')
plt.title("Visualizing Important Features")
plt.legend()
plt.show()

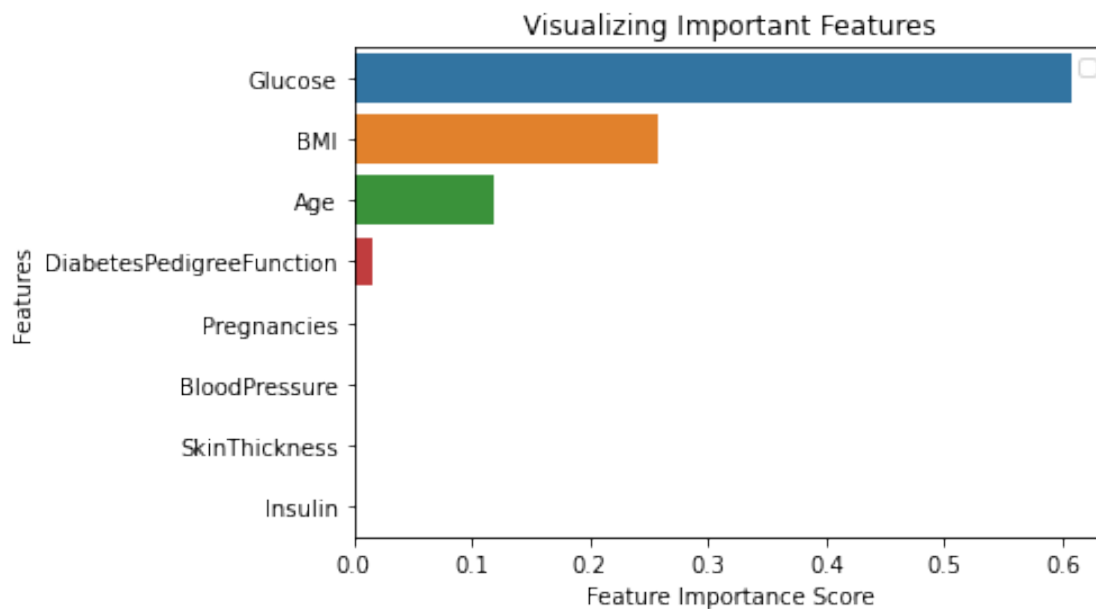
```

```

Glucose          0.606802
BMI              0.258369
Age              0.118413
DiabetesPedigreeFunction  0.016416
Pregnancies      0.000000
BloodPressure    0.000000
SkinThickness    0.000000
Insulin          0.000000
dtype: float64

```

No handles with labels found to put in legend.



2.0.8 criterion="entropy", splitter="random", max_depth=3

```

[61]: # Define and build model
      clf = DecisionTreeClassifier(criterion="entropy", splitter="random",
      ↪max_depth=3)

```

```
clf = clf.fit(x_train,y_train)
y_pred = clf.predict(x_test)
```

```
[62]: print(y_pred)
```

```
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 1 1 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 1 0 0]
```

```
[63]: print(np.array(y_test))
```

```
[0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1 1 0 0 0 1 1 1 1 0 0 0 1 0 1 1 0 0 1 0 1 0
 0 0 0 0 0 0 0 1 0 0 1 1 0 1 0 0 1 0 1 0 1 0 0 0 0 0 1 0 1 0 1 1 0 1 1 0 0 0
 1 1 1 0 0 1 1 0 1 1 0 0 1 0 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 1 0 1 0 1
 0 0 0 1 0 0 1 0 1 0 1 1 0 0 0 0 1 0 0 1 0 1 0 0 0 0 0 1 0 1 0 0 1 1 1 0 0
 1 0 0 1 0 0]
```

```
[64]: accuracy["dt_entropy_random_3"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

```
Accuracy: 0.6818181818181818
```

```
[65]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[98  1]
 [48  7]]
```

```
[66]: print(metrics.classification_report(y_test, y_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.67 | 0.99 | 0.80 | 99 |
| 1 | 0.88 | 0.13 | 0.22 | 55 |
| accuracy | | | 0.68 | 154 |
| macro avg | 0.77 | 0.56 | 0.51 | 154 |
| weighted avg | 0.74 | 0.68 | 0.59 | 154 |

3 Accuracy visulization of Decision Tree

```
[67]: accuracy_df_dt = pd.DataFrame(list(zip(accuracy.keys(), accuracy.values()))),
    columns = ['Arguments', 'Accuracy'])
accuracy_df_dt
```

```
[67]:      Arguments  Accuracy
0      dt_gini_best  0.701299
```

```

1      dt_gini_best_8  0.668831
2      dt_entropy_best 0.668831
3      dt_entropy_best_8 0.707792
4      dt_entropy_random 0.623377
5      dt_entropy_random_8 0.707792
6      dt_entropy_best_3 0.792208
7      dt_entropy_random_3 0.681818

```

```
[68]: fig = px.bar(accuracy_df_dt, x='Arguments', y='Accuracy')
fig.show()
```

4 Random Forest

```
[69]: accuracy_rf = {}
```

4.0.1 n_estimators = 1000, criterion='entropy'

```
[70]: # Instantiate model with 1000 decision trees
rf = RandomForestClassifier(n_estimators = 1000, criterion='entropy')
# Train the model on training data
rf.fit(x_train,y_train)
# Use the forest's predict method on the test data
y_pred = rf.predict(x_test)
```

```
[71]: print(y_pred)
```

```

[1 0 0 0 0 0 0 0 0 0 0 1 0 1 1 0 1 0 0 0 0 1 0 1 0 0 0 0 1 0 1 0 0 0 1 0 1 0
 0 0 0 0 0 0 1 0 0 1 1 0 0 0 0 1 0 1 0 0 0 0 0 1 0 1 0 1 0 0 1 1 1 1 1 0 0
 1 0 1 0 0 1 1 0 0 0 0 1 0 0 1 0 0 0 0 0 1 1 0 0 1 0 0 0 1 0 0 0 1 1 0 0 0
 0 0 0 0 0 0 0 0 1 0 1 0 0 0 0 0 0 1 0 1 0 1 0 0 0 0 1 0 0 1 0 0 1 1 1 0 0
 0 0 0 1 1 0]

```

```
[72]: print(np.array(y_test))
```

```

[0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1 1 0 0 0 1 1 1 1 0 0 0 1 0 1 1 0 0 1 0 1 0
 0 0 0 0 0 0 1 0 0 1 1 0 1 0 0 1 0 1 0 1 0 0 0 0 0 1 0 1 0 1 1 0 1 1 0 0 0
 1 1 1 0 0 1 1 0 1 1 0 0 1 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 1 0 1 0 1
 0 0 0 1 0 0 1 0 1 0 1 1 0 0 0 0 1 0 0 1 0 1 0 0 0 0 0 1 0 1 0 0 1 1 1 0 0
 1 0 0 1 0 0]

```

```
[73]: accuracy_rf["rf_entropy_1000"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.7987012987012987

```
[74]: print(metrics.confusion_matrix(y_test, y_pred))
```

```

[[87 12]
 [19 36]]

```


```
[75]: print(metrics.classification_report(y_test, y_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.82 | 0.88 | 0.85 | 99 |
| 1 | 0.75 | 0.65 | 0.70 | 55 |
| accuracy | | | 0.80 | 154 |
| macro avg | 0.79 | 0.77 | 0.77 | 154 |
| weighted avg | 0.80 | 0.80 | 0.80 | 154 |

```
[76]: feature_imp = pd.Series(rf.feature_importances_,index=X.columns).
      ↪sort_values(ascending=False)
print(feature_imp)
# Creating a bar plot
sns.barplot(x=feature_imp, y=feature_imp.index)
# Add labels to your graph
plt.xlabel('Feature Importance Score')
plt.ylabel('Features')
plt.title("Visualizing Important Features")
plt.legend()
plt.show()
```

No handles with labels found to put in legend.

| | |
|--------------------------|----------|
| Glucose | 0.237764 |
| BMI | 0.172892 |
| Age | 0.131403 |
| DiabetesPedigreeFunction | 0.128862 |
| BloodPressure | 0.092695 |
| Pregnancies | 0.082019 |
| Insulin | 0.078341 |
| SkinThickness | 0.076023 |
| dtype: | float64 |



4.0.2 n_estimators = 100, criterion='entropy'

```
[77]: # Instantiate model with 100 decision trees
rf = RandomForestClassifier(n_estimators = 100, criterion='entropy')
# Train the model on training data
rf.fit(x_train,y_train)
# Use the forest's predict method on the test data
y_pred = rf.predict(x_test)
```

```
[78]: print(y_pred)
```

```
[1 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1 0 0 0 0 1 0 1 0 0 0 0 1 0 1 0 0 0 1 0 1 0
0 0 0 0 0 0 0 0 0 0 1 1 0 0 0 0 1 0 1 0 0 0 0 0 1 0 1 0 1 0 0 0 1 1 1 1 0 0
1 1 1 0 0 1 1 0 1 0 0 1 0 0 1 0 0 0 0 0 0 1 1 0 0 1 0 0 0 1 0 0 0 1 1 0 0 0
0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 1 0 1 0 1 0 0 0 0 0 1 0 0 1 0 0 1 1 0 0 0
0 0 0 1 1 0]
```

```
[79]: print(np.array(y_test))
```

```
[0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1 1 0 0 0 1 1 1 1 0 0 0 1 0 1 1 0 0 1 0 1 0
0 0 0 0 0 0 1 0 0 1 1 0 1 0 0 1 0 1 0 1 0 0 0 0 0 1 0 1 0 1 1 0 1 1 0 0 0
1 1 1 0 0 1 1 0 1 1 0 0 1 0 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 1 0 1 0 1
0 0 0 1 0 0 1 0 1 0 1 1 0 0 0 0 1 0 0 1 0 1 0 0 0 0 0 1 0 1 0 0 1 1 1 0 0
1 0 0 1 0 0]
```

```
[80]: accuracy_rf["rf_entropy_100"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.7922077922077922

```
[81]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[88 11]
 [21 34]]
```

```
[82]: print(metrics.classification_report(y_test, y_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.81 | 0.89 | 0.85 | 99 |
| 1 | 0.76 | 0.62 | 0.68 | 55 |
| accuracy | | | 0.79 | 154 |
| macro avg | 0.78 | 0.75 | 0.76 | 154 |
| weighted avg | 0.79 | 0.79 | 0.79 | 154 |

4.0.3 n_estimators = 1000, random_state = 42, criterion='entropy'

```
[83]: # Instantiate model with 1000 decision trees
rf = RandomForestClassifier(n_estimators = 1000, random_state = 42,
    ↪criterion='entropy')
# Train the model on training data
rf.fit(x_train,y_train)
# Use the forest's predict method on the test data
y_pred = rf.predict(x_test)
```

```
[84]: print(y_pred)
```

```
[1 0 0 0 0 0 0 0 0 0 0 1 0 1 1 0 1 0 0 0 0 1 0 1 0 0 0 0 1 0 1 0 0 0 1 0 1 0
 0 0 0 0 0 0 1 0 0 1 1 0 0 0 0 1 0 1 0 0 0 0 0 1 0 1 0 1 0 0 1 1 1 1 1 1 0
 1 1 1 0 0 1 1 0 0 0 0 1 1 0 1 0 0 0 0 0 1 1 0 0 1 0 0 0 1 0 0 0 1 1 0 0 0
 0 0 0 0 0 0 0 0 1 0 1 0 0 0 0 0 0 0 1 0 1 0 1 0 0 0 0 1 0 0 1 0 0 1 1 0 0 0
 0 0 0 1 1 0]
```

```
[85]: print(np.array(y_test))
```

```
[0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1 1 0 0 0 1 1 1 1 0 0 0 1 0 1 1 0 0 1 0 1 0
 0 0 0 0 0 0 1 0 0 1 1 0 1 0 0 1 0 1 0 1 0 0 0 0 0 1 0 1 0 1 1 0 1 1 0 0 0
 1 1 1 0 0 1 1 0 1 1 0 0 1 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 1 0 1 0 1
 0 0 0 1 0 0 1 0 1 0 1 1 0 0 0 0 1 0 0 1 0 1 0 0 0 0 0 1 0 1 0 0 1 1 1 0 0
 1 0 0 1 0 0]
```

```
[86]: accuracy_rf["rf_entropy_1000_42"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.7987012987012987

```
[87]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[86 13]
 [18 37]]
```

```
[88]: print(metrics.classification_report(y_test, y_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.83 | 0.87 | 0.85 | 99 |
| 1 | 0.74 | 0.67 | 0.70 | 55 |
| accuracy | | | 0.80 | 154 |
| macro avg | 0.78 | 0.77 | 0.78 | 154 |
| weighted avg | 0.80 | 0.80 | 0.80 | 154 |

4.0.4 n_estimators = 100, random_state = 42, criterion='entropy'

```
[89]: # Instantiate model with 100 decision trees
rf = RandomForestClassifier(n_estimators = 100, random_state = 42, max_depth = 8, criterion='entropy')
# Train the model on training data
rf.fit(x_train,y_train)
# Use the forest's predict method on the test data
y_pred = rf.predict(x_test)
```

```
[90]: print(y_pred)
```

```
[0 0 0 0 0 0 0 0 0 0 0 1 0 1 1 0 1 0 0 0 0 1 0 1 0 0 1 0 1 0 1 0 0 0 1 0 1 0
 0 0 0 0 0 0 0 0 0 0 1 1 0 0 0 0 1 0 1 0 0 0 0 0 1 0 1 0 1 0 0 0 1 1 1 1 0 0
 1 0 1 0 0 1 1 0 0 0 0 1 1 0 1 0 0 0 0 0 1 1 0 0 1 0 0 0 1 0 0 0 1 1 0 0 1
 0 1 0 0 0 0 0 0 1 0 1 0 0 0 0 0 0 1 0 1 0 1 0 0 0 0 1 0 0 1 0 0 1 1 0 0 0
 0 0 0 1 1 0]
```

```
[91]: print(np.array(y_test))
```

```
[0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1 1 0 0 0 1 1 1 1 0 0 0 1 0 1 1 0 0 1 0 1 0
 0 0 0 0 0 0 1 0 0 1 1 0 1 0 0 1 0 1 0 1 0 0 0 0 0 1 0 1 0 1 1 0 1 1 0 0 0
 1 1 1 0 0 1 1 0 1 1 0 0 1 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 1 0 1 0 1
 0 0 0 1 0 0 1 0 1 0 1 1 0 0 0 0 1 0 0 1 0 1 0 0 0 0 0 1 0 1 0 0 1 1 1 0 0
 1 0 0 1 0 0]
```

```
[92]: accuracy_rf["rf_entropy_100_42"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

```
Accuracy: 0.7857142857142857
```

```
[93]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[86 13]
 [20 35]]
```

```
[94]: print(metrics.classification_report(y_test, y_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.81 | 0.87 | 0.84 | 99 |
| 1 | 0.73 | 0.64 | 0.68 | 55 |
| accuracy | | | 0.79 | 154 |
| macro avg | 0.77 | 0.75 | 0.76 | 154 |
| weighted avg | 0.78 | 0.79 | 0.78 | 154 |

4.0.5 `n_estimators = 1000, random_state = 42, max_depth = 8, criterion='entropy'`

```
[95]: # Instantiate model with 1000 decision trees
rf = RandomForestClassifier(n_estimators = 1000, random_state = 42, max_depth = 8, criterion='entropy')
# Train the model on training data
rf.fit(x_train,y_train)
# Use the forest's predict method on the test data
y_pred = rf.predict(x_test)
```

```
[96]: print(y_pred)
```

```
[1 0 0 0 0 0 0 0 0 0 0 1 0 1 1 0 1 0 0 0 0 1 0 1 0 0 0 0 1 0 1 0 0 0 1 0 1 0
0 0 0 0 0 0 0 0 0 0 1 1 0 0 0 0 1 0 1 0 0 0 0 0 1 0 1 0 1 0 0 0 1 1 1 1 0 0
1 0 1 0 0 1 1 0 0 0 0 1 1 0 1 0 0 0 0 0 1 1 0 0 1 0 0 0 1 0 0 0 1 1 0 0 0
0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 1 0 1 0 1 0 0 0 0 1 0 0 1 0 0 1 1 1 0 0
0 0 0 1 1 0]
```

```
[97]: print(np.array(y_test))
```

```
[0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1 1 0 0 0 1 1 1 1 0 0 0 1 0 1 1 0 0 1 0 1 0
0 0 0 0 0 0 1 0 0 1 1 0 1 0 0 1 0 1 0 1 0 0 0 0 0 1 0 1 0 1 1 0 1 1 0 0 0
1 1 1 0 0 1 1 0 1 1 0 0 1 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 1 0 1 0 1
0 0 0 1 0 0 1 0 1 0 1 1 0 0 0 0 1 0 0 1 0 1 0 0 0 0 0 1 0 1 0 0 1 1 1 0 0
1 0 0 1 0 0]
```

```
[98]: accuracy_rf["rf_entropy_1000_42_8"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:", metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.7857142857142857

```
[99]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[87 12]
 [21 34]]
```

```
[100]: print(metrics.classification_report(y_test, y_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.81 | 0.88 | 0.84 | 99 |
| 1 | 0.74 | 0.62 | 0.67 | 55 |
| accuracy | | | 0.79 | 154 |
| macro avg | 0.77 | 0.75 | 0.76 | 154 |
| weighted avg | 0.78 | 0.79 | 0.78 | 154 |

4.0.6 n_estimators = 100, random_state = 42, max_depth = 8, criterion='entropy'

```
[101]: # Instantiate model with 100 decision trees
rf = RandomForestClassifier(n_estimators = 100, random_state = 42, max_depth = 8, criterion='entropy')
# Train the model on training data
rf.fit(x_train,y_train)
# Use the forest's predict method on the test data
y_pred = rf.predict(x_test)
```

```
[102]: print(y_pred)
```

```
[0 0 0 0 0 0 0 0 0 0 1 0 1 1 0 1 0 0 0 0 1 0 1 0 0 1 0 1 0 1 0 0 0 1 0 1 0
0 0 0 0 0 0 0 0 0 0 1 1 0 0 0 0 1 0 1 0 0 0 0 0 1 0 1 0 1 0 0 0 1 1 1 1 0 0
1 0 1 0 0 1 1 0 0 0 0 1 1 0 1 0 0 0 0 0 1 1 0 0 1 0 0 0 1 0 0 0 1 1 0 0 1
0 1 0 0 0 0 0 0 1 0 1 0 0 0 0 0 0 1 0 1 0 1 0 0 0 0 1 0 0 1 0 0 1 1 0 0 0
0 0 0 1 1 0]
```

```
[103]: print(np.array(y_test))
```

```
[0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1 1 0 0 0 1 1 1 1 0 0 0 1 0 1 1 0 0 1 0 1 0
0 0 0 0 0 0 1 0 0 1 1 0 1 0 0 1 0 1 0 1 0 0 0 0 0 1 0 1 0 1 1 0 1 1 0 0 0
1 1 1 0 0 1 1 0 1 1 0 0 1 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 1 0 1 0 1
0 0 0 1 0 0 1 0 1 0 1 1 0 0 0 0 1 0 0 1 0 1 0 0 0 0 0 1 0 1 0 0 1 1 1 0 0
1 0 0 1 0 0]
```

```
[104]: accuracy_rf["rf_entropy_100_42_8"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.7857142857142857

```
[105]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[86 13]
 [20 35]]
```

```
[106]: print(metrics.classification_report(y_test, y_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.81 | 0.87 | 0.84 | 99 |
| 1 | 0.73 | 0.64 | 0.68 | 55 |
| accuracy | | | 0.79 | 154 |
| macro avg | 0.77 | 0.75 | 0.76 | 154 |
| weighted avg | 0.78 | 0.79 | 0.78 | 154 |

4.0.7 n_estimators = 1000

```
[107]: # Instantiate model with 1000 decision trees
rf = RandomForestClassifier(n_estimators = 1000)
# Train the model on training data
rf.fit(x_train,y_train)
# Use the forest's predict method on the test data
y_pred = rf.predict(x_test)
```

```
[108]: print(y_pred)
```

```
[1 0 0 0 0 0 0 0 0 0 0 1 0 1 1 0 1 0 0 0 0 1 0 1 0 0 0 0 1 0 1 0 0 0 1 0 1 0
0 0 0 0 0 0 0 0 0 0 1 1 0 0 0 0 1 0 1 0 0 0 0 0 1 0 1 0 1 0 0 0 1 1 1 1 0 0
1 1 1 0 0 1 1 0 0 0 0 1 1 0 1 0 0 0 0 0 1 1 0 0 1 0 0 0 1 0 0 0 1 1 0 0 1
0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 1 0 1 0 1 0 0 0 0 1 0 0 1 0 0 1 1 1 0 0
0 0 0 1 1 0]
```

```
[109]: print(np.array(y_test))
```

```
[0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1 1 0 0 0 1 1 1 1 0 0 0 1 0 1 1 0 0 1 0 1 0
0 0 0 0 0 0 1 0 0 1 1 0 1 0 0 1 0 1 0 1 0 0 0 0 0 1 0 1 0 1 1 0 1 1 0 0 0
1 1 1 0 0 1 1 0 1 1 0 0 1 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 1 0 1 0 1
0 0 0 1 0 0 1 0 1 0 1 1 0 0 0 0 1 0 0 1 0 1 0 0 0 0 0 1 0 1 0 0 1 1 1 0 0
1 0 0 1 0 0]
```

```
[110]: accuracy_rf["rf_gini_1000"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.7987012987012987

```
[111]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[87 12]
 [19 36]]
```

```
[112]: print(metrics.classification_report(y_test, y_pred))
```

| | precision | recall | f1-score | support |
|---|-----------|--------|----------|---------|
| 0 | 0.82 | 0.88 | 0.85 | 99 |

| | | | | | |
|--------------|------|------|------|------|-----|
| | 1 | 0.75 | 0.65 | 0.70 | 55 |
| accuracy | | | | 0.80 | 154 |
| macro avg | 0.79 | 0.77 | 0.77 | | 154 |
| weighted avg | 0.80 | 0.80 | 0.80 | | 154 |

4.0.8 n_estimators = 100

```
[113]: # Instantiate model with 100 decision trees
rf = RandomForestClassifier(n_estimators = 100)
# Train the model on training data
rf.fit(x_train,y_train)
# Use the forest's predict method on the test data
y_pred = rf.predict(x_test)
```

```
[114]: print(y_pred)
```

```
[1 0 0 0 0 0 0 0 0 0 1 0 1 1 0 1 0 1 0 0 0 0 1 0 0 0 0 1 0 1 0 0 0 1 0 1 0
0 0 0 0 0 0 1 0 0 1 1 0 0 0 0 1 0 1 0 0 0 0 0 1 0 1 0 1 0 0 1 1 1 1 1 1 0
1 1 1 0 0 1 1 0 0 0 0 1 1 0 1 0 0 0 0 0 1 1 0 0 1 0 0 0 1 0 0 0 1 1 0 0 1
0 1 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 1 0 1 0 1 0 0 0 0 1 0 0 1 0 0 1 0 0 0 0
0 0 0 1 1 0]
```

```
[115]: print(np.array(y_test))
```

```
[0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1 1 0 0 0 1 1 1 1 0 0 0 1 0 1 1 0 0 1 0 1 0
0 0 0 0 0 0 1 0 0 1 1 0 1 0 0 1 0 1 0 1 0 0 0 0 0 1 0 1 0 1 1 0 1 1 0 0 0
1 1 1 0 0 1 1 0 1 1 0 0 1 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 1 0 1 0 1
0 0 0 1 0 0 1 0 1 0 1 1 0 0 0 0 1 0 0 1 0 1 0 0 0 0 0 1 0 1 0 0 1 1 1 0 0
1 0 0 1 0 0]
```

```
[116]: accuracy_rf["rf_gini_100"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.7727272727272727

```
[117]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[84 15]
 [20 35]]
```

```
[118]: print(metrics.classification_report(y_test, y_pred))
```

| | precision | recall | f1-score | support |
|----------|-----------|--------|----------|---------|
| 0 | 0.81 | 0.85 | 0.83 | 99 |
| 1 | 0.70 | 0.64 | 0.67 | 55 |
| accuracy | | | 0.77 | 154 |

| | | | | |
|--------------|------|------|------|-----|
| macro avg | 0.75 | 0.74 | 0.75 | 154 |
| weighted avg | 0.77 | 0.77 | 0.77 | 154 |

4.0.9 n_estimators = 1000, random_state = 42

```
[119]: # Instantiate model with 1000 decision trees
rf = RandomForestClassifier(n_estimators = 1000, random_state = 42)
# Train the model on training data
rf.fit(x_train,y_train)
# Use the forest's predict method on the test data
y_pred = rf.predict(x_test)
```

```
[120]: print(y_pred)

[0 0 0 0 0 0 0 0 0 0 0 1 0 1 1 0 1 0 0 0 0 1 0 1 0 0 0 0 1 0 1 0 0 0 1 0 1 0
 0 0 0 0 0 0 0 0 0 0 1 1 0 0 0 0 1 0 1 0 0 0 0 0 1 0 1 0 1 0 0 0 1 1 1 1 0 0
 1 1 1 0 0 1 1 0 0 0 0 1 1 0 1 0 0 0 0 0 1 1 0 0 1 0 0 0 1 0 0 0 1 1 0 0 1
 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 1 0 1 0 1 0 0 0 0 1 0 0 1 0 0 1 1 0 0 0
 0 0 0 1 1 0]
```

```
[121]: print(np.array(y_test))

[0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1 1 0 0 0 1 1 1 1 0 0 0 1 0 1 1 0 0 1 0 1 0
 0 0 0 0 0 0 1 0 0 1 1 0 1 0 0 1 0 1 0 1 0 0 0 0 0 1 0 1 0 1 1 0 1 1 0 0 0
 1 1 1 0 0 1 1 0 1 1 0 0 1 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 1 0 1 0 1
 0 0 0 1 0 0 1 0 1 0 1 1 0 0 0 0 1 0 0 1 0 1 0 0 0 0 0 1 0 1 0 0 1 1 1 0 0
 1 0 0 1 0 0]
```

```
[122]: accuracy_rf["rf_gini_1000_42"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.7987012987012987

```
[123]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[88 11]
 [20 35]]
```

```
[124]: print(metrics.classification_report(y_test, y_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.81 | 0.89 | 0.85 | 99 |
| 1 | 0.76 | 0.64 | 0.69 | 55 |
| accuracy | | | 0.80 | 154 |
| macro avg | 0.79 | 0.76 | 0.77 | 154 |
| weighted avg | 0.80 | 0.80 | 0.79 | 154 |

4.0.10 n_estimators = 100, random_state = 42

```
[125]: # Instantiate model with 100 decision trees
rf = RandomForestClassifier(n_estimators = 100, random_state = 42, max_depth = 8)
# Train the model on training data
rf.fit(x_train,y_train)
# Use the forest's predict method on the test data
y_pred = rf.predict(x_test)
```

```
[126]: print(y_pred)
```

```
[0 0 0 0 0 0 0 0 0 0 0 1 0 1 1 0 1 0 0 0 0 1 0 1 0 0 1 0 1 0 1 0 0 0 1 0 1 0
 0 0 1 0 0 0 1 0 0 1 1 0 0 0 0 1 0 1 0 0 0 0 0 1 0 1 0 1 0 0 0 1 1 1 1 0 0
 1 0 1 0 0 1 1 0 0 0 0 1 1 0 1 0 0 0 0 0 1 1 1 0 1 0 0 0 1 0 0 0 1 1 0 0 1
 0 0 0 0 0 0 1 0 1 0 1 0 0 0 0 0 0 1 0 1 0 1 0 0 0 0 1 0 0 1 0 0 1 1 0 0 0
 0 0 0 1 1 0]
```

```
[127]: print(np.array(y_test))
```

```
[0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1 1 0 0 0 1 1 1 1 0 0 0 1 0 1 1 0 0 1 0 1 0
 0 0 0 0 0 0 1 0 0 1 1 0 1 0 0 1 0 1 0 1 0 0 0 0 0 1 0 1 0 1 1 0 1 1 0 0 0
 1 1 1 0 0 1 1 0 1 1 0 0 1 0 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 1 0 1 0 1
 0 0 0 1 0 0 1 0 1 0 1 1 0 0 0 0 1 0 0 1 0 1 0 0 0 0 0 1 0 1 0 0 1 1 1 0 0
 1 0 0 1 0 0]
```

```
[128]: accuracy_rf["rf_gini_100_42"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.7922077922077922

```
[129]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[85 14]
 [18 37]]
```

```
[130]: print(metrics.classification_report(y_test, y_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.83 | 0.86 | 0.84 | 99 |
| 1 | 0.73 | 0.67 | 0.70 | 55 |
| accuracy | | | 0.79 | 154 |
| macro avg | 0.78 | 0.77 | 0.77 | 154 |
| weighted avg | 0.79 | 0.79 | 0.79 | 154 |

4.0.11 n_estimators = 1000, random_state = 42, max_depth = 8

```
[131]: # Instantiate model with 1000 decision trees
rf = RandomForestClassifier(n_estimators = 1000, random_state = 42, max_depth = 8)
# Train the model on training data
rf.fit(x_train,y_train)
# Use the forest's predict method on the test data
y_pred = rf.predict(x_test)
```

```
[132]: print(y_pred)

[1 0 0 0 0 0 0 0 0 0 0 1 0 1 1 0 1 0 0 0 0 1 0 1 0 0 0 0 1 0 1 0 0 0 1 0 1 0
 0 0 0 0 0 0 0 0 0 0 1 1 0 0 0 0 1 0 1 0 0 0 0 0 1 0 1 0 1 0 0 0 1 1 1 1 0 0
 1 0 1 0 0 1 1 0 0 0 0 1 1 0 1 0 0 0 0 0 1 1 0 0 1 0 0 0 1 0 0 0 1 1 0 0 1
 0 1 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 1 0 1 0 1 0 0 0 0 1 0 0 1 0 0 1 1 0 0 0
 0 0 0 1 1 0]
```

```
[133]: print(np.array(y_test))

[0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1 1 0 0 0 1 1 1 1 0 0 0 1 0 1 1 0 0 1 0 1 0
 0 0 0 0 0 0 1 0 0 1 1 0 1 0 0 1 0 1 0 1 0 0 0 0 0 1 0 1 0 1 1 0 1 1 0 0 0
 1 1 1 0 0 1 1 0 1 1 0 0 1 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 1 0 1 0 1
 0 0 0 1 0 0 1 0 1 0 1 1 0 0 0 0 1 0 0 1 0 1 0 0 0 0 0 1 0 1 0 0 1 1 1 0 0
 1 0 0 1 0 0]
```

```
[134]: accuracy_rf["rf_gini_1000_42_8"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.7792207792207793

```
[135]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[86 13]
 [21 34]]
```

```
[136]: print(metrics.classification_report(y_test, y_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.80 | 0.87 | 0.83 | 99 |
| 1 | 0.72 | 0.62 | 0.67 | 55 |
| accuracy | | | 0.78 | 154 |
| macro avg | 0.76 | 0.74 | 0.75 | 154 |
| weighted avg | 0.78 | 0.78 | 0.77 | 154 |

4.0.12 n_estimators = 100, random_state = 42, max_depth = 8

```
[137]: # Instantiate model with 100 decision trees
rf = RandomForestClassifier(n_estimators = 100, random_state = 42, max_depth = 8)
# Train the model on training data
rf.fit(x_train,y_train)
# Use the forest's predict method on the test data
y_pred = rf.predict(x_test)
```

```
[138]: print(y_pred)
```

```
[0 0 0 0 0 0 0 0 0 0 1 0 1 1 0 1 0 0 0 0 1 0 1 0 0 1 0 1 0 1 0 0 0 1 0 1 0
 0 0 1 0 0 0 1 0 0 1 1 0 0 0 0 1 0 1 0 0 0 0 0 1 0 1 0 1 0 0 0 1 1 1 1 0 0
 1 0 1 0 0 1 1 0 0 0 0 1 1 0 1 0 0 0 0 0 1 1 1 0 1 0 0 0 1 0 0 0 1 1 0 0 1
 0 0 0 0 0 0 1 0 1 0 1 0 0 0 0 0 0 1 0 1 0 1 0 0 0 0 1 0 0 1 0 0 1 1 0 0 0
 0 0 0 1 1 0]
```

```
[139]: print(np.array(y_test))
```

```
[0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1 1 0 0 0 1 1 1 1 0 0 0 1 0 1 1 0 0 1 0 1 0
 0 0 0 0 0 0 1 0 0 1 1 0 1 0 0 1 0 1 0 1 0 0 0 0 0 1 0 1 0 1 1 0 1 1 0 0 0
 1 1 1 0 0 1 1 0 1 1 0 0 1 0 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 1 0 1 0 1
 0 0 0 1 0 0 1 0 1 0 1 1 0 0 0 0 1 0 0 1 0 1 0 0 0 0 0 1 0 1 0 0 1 1 1 0 0
 1 0 0 1 0 0]
```

```
[140]: accuracy_rf["rf_gini_100_42_8"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.7922077922077922

```
[141]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[85 14]
 [18 37]]
```

```
[142]: print(metrics.classification_report(y_test, y_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.83 | 0.86 | 0.84 | 99 |
| 1 | 0.73 | 0.67 | 0.70 | 55 |
| accuracy | | | 0.79 | 154 |
| macro avg | 0.78 | 0.77 | 0.77 | 154 |
| weighted avg | 0.79 | 0.79 | 0.79 | 154 |

5 Accuracy visulization of Random Forest

```
[143]: accuracy_df_rf = pd.DataFrame(list(zip(accuracy_rf.keys(), accuracy_rf.  
    ↪values()))), columns=['Arguments', 'Accuracy'])  
accuracy_df_rf
```

```
[143]:
```

| | Arguments | Accuracy |
|----|----------------------|----------|
| 0 | rf_entropy_1000 | 0.798701 |
| 1 | rf_entropy_100 | 0.792208 |
| 2 | rf_entropy_1000_42 | 0.798701 |
| 3 | rf_entropy_100_42 | 0.785714 |
| 4 | rf_entropy_1000_42_8 | 0.785714 |
| 5 | rf_entropy_100_42_8 | 0.785714 |
| 6 | rf_gini_1000 | 0.798701 |
| 7 | rf_gini_100 | 0.772727 |
| 8 | rf_gini_1000_42 | 0.798701 |
| 9 | rf_gini_100_42 | 0.792208 |
| 10 | rf_gini_1000_42_8 | 0.779221 |
| 11 | rf_gini_100_42_8 | 0.792208 |

```
[144]: fig = px.bar(accuracy_df_rf, x='Arguments', y='Accuracy')  
fig.show()
```

```
[145]: accuracy_df = pd.concat([accuracy_df_dt, accuracy_df_rf])  
accuracy_df['Accuracy'] = round(accuracy_df['Accuracy'] * 100, 2)  
fig = px.bar(accuracy_df, x='Arguments', y='Accuracy')  
print(accuracy_df['Accuracy'].max())  
fig.show()
```

79.87