

Decision Tree & Random Forest V6

November 21, 2021

Replace all zero features with mean
RandomUnderSampler

```
[1]: import numpy as np # Import numpy for data preprocessing
import pandas as pd # Import pandas for data frame read
import matplotlib.pyplot as plt # Import matplotlib for data visualisation
import seaborn as sns # Import seaborn for data visualisation
import plotly.express as px # Import plotly for data visualisation
from sklearn.model_selection import train_test_split # Import train_test_split
    ↳ for data split
from sklearn.tree import DecisionTreeClassifier # Import Decision Tree
    ↳ Classifier
from sklearn.ensemble import RandomForestClassifier # Import Random Forest
    ↳ Classifier
from sklearn.model_selection import train_test_split # Import train_test_split
    ↳ function
from sklearn import metrics # Import scikit-learn metrics module for accuracy
    ↳ calculation
from sklearn import tree # Import export_graphviz for visualizing Decision Trees
from imblearn.under_sampling import RandomUnderSampler # Up-sample or
    ↳ Down-sample
```

0.1 Data read

```
[2]: df = pd.read_csv("data/diabetes.csv") # Data read
```

```
[3]: df.head() # print data
```

```
[3]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0

2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[4]: df.isna().sum() # check for null value
```

```
[4]: Pregnancies      0
      Glucose          0
      BloodPressure    0
      SkinThickness    0
      Insulin          0
      BMI              0
      DiabetesPedigreeFunction  0
      Age              0
      Outcome          0
      dtype: int64
```

```
[5]: df.describe()
```

```
[5]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479
std	3.369578	31.972618	19.355807	15.952218	115.244002
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

```
[6]: df.corr()
```

```
[6]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness \
Pregnancies	1.000000	0.129459	0.141282	-0.081672
Glucose	0.129459	1.000000	0.152590	0.057328
BloodPressure	0.141282	0.152590	1.000000	0.207371
SkinThickness	-0.081672	0.057328	0.207371	1.000000
Insulin	-0.073535	0.331357	0.088933	0.436783


BMI	0.017683	0.221071	0.281805	0.392573
DiabetesPedigreeFunction	-0.033523	0.137337	0.041265	0.183928
Age	0.544341	0.263514	0.239528	-0.113970
Outcome	0.221898	0.466581	0.065068	0.074752

	Insulin	BMI	DiabetesPedigreeFunction	\
Pregnancies	-0.073535	0.017683	-0.033523	
Glucose	0.331357	0.221071	0.137337	
BloodPressure	0.088933	0.281805	0.041265	
SkinThickness	0.436783	0.392573	0.183928	
Insulin	1.000000	0.197859	0.185071	
BMI	0.197859	1.000000	0.140647	
DiabetesPedigreeFunction	0.185071	0.140647	1.000000	
Age	-0.042163	0.036242	0.033561	
Outcome	0.130548	0.292695	0.173844	

	Age	Outcome
Pregnancies	0.544341	0.221898
Glucose	0.263514	0.466581
BloodPressure	0.239528	0.065068
SkinThickness	-0.113970	0.074752
Insulin	-0.042163	0.130548
BMI	0.036242	0.292695
DiabetesPedigreeFunction	0.033561	0.173844
Age	1.000000	0.238356
Outcome	0.238356	1.000000

```
[7]: sns.heatmap(df.corr())
```

```
[7]: <AxesSubplot:>
```



1 Data split

```
[8]: df.shape
```

```
[8]: (768, 9)
```

```
[9]: X = df.iloc[:,0:-1] # All features
      Y = df.iloc[:, -1] # Target
```

```
[10]: X.head()
```

```
[10]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age
0	0.627	50
1	0.351	31
2	0.672	32
3	0.167	21
4	2.288	33

```
[11]: Y.head()
```

```
[11]: 0    1
      1    0
      2    1
      3    0
      4    1
      Name: Outcome, dtype: int64
```

```
[12]: print("X.shape : ", X.shape)
      print("Y.shape : ", Y.shape)
```

```
X.shape : (768, 8)
Y.shape : (768,)
```

```
[13]: rus = RandomUnderSampler(random_state=42)
      X_res, Y_res = rus.fit_resample(X, Y)
```

```
[14]: print("X_res.shape : ", X_res.shape)
      print("Y_res.shape : ", Y_res.shape)
```

```
X_res.shape : (536, 8)
Y_res.shape : (536,)
```

```
[15]: # Data split
      x_train, x_test, y_train, y_test = train_test_split(X_res, Y_res, test_size=0.
      ↪2, random_state=1)
      # x_dev, x_test, y_dev, y_test = train_test_split(x_test, y_test, test_size= 0.
      ↪5)
```

```
[16]: print("Original data size : ", X.shape, Y.shape)
      print("Train data size : ", x_train.shape, y_train.shape)
      # print("Dev data size : ", x_dev.shape, y_dev.shape)
      print("Test data size : ", x_test.shape, y_test.shape)
```

```
Original data size : (768, 8) (768,)
Train data size : (428, 8) (428,)
Test data size : (108, 8) (108,)
```

2 Preprocessing

```
[17]: # replace zero bmi value with it's mean
print("Before BMI mean : ",round(x_train.loc[:, 'BMI'].mean(),1))
x_test.loc[:, 'BMI'] = x_test.loc[:, 'BMI'].replace(0, x_train.loc[:, 'BMI'].
    ↪mean())
x_train.loc[:, 'BMI'] = x_train.loc[:, 'BMI'].replace(0, x_train.loc[:, 'BMI'].
    ↪mean())
print("After BMI mean : ",round(x_train.loc[:, 'BMI'].mean(),1))
```

Before BMI mean : 32.4

After BMI mean : 33.0

/Users/kamal/opt/anaconda3/lib/python3.8/site-
packages/pandas/core/indexing.py:1773: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
self._setitem_single_column(ilocs[0], value, pi)

```
[18]: # replace zero skinthickness value with it's mean
print("Before SkinThickness mean : ",round(x_train.loc[:, 'SkinThickness'].
    ↪mean(),1))
x_test.loc[:, 'SkinThickness'] = x_test.loc[:, 'SkinThickness'].replace(0,
    ↪x_train.loc[:, 'SkinThickness'].mean())
x_train.loc[:, 'SkinThickness'] = x_train.loc[:, 'SkinThickness'].replace(0,
    ↪x_train.loc[:, 'SkinThickness'].mean())
print("After SkinThickness mean : ",round(x_train.loc[:, 'SkinThickness'].
    ↪mean(),1))
```

Before SkinThickness mean : 20.7

After SkinThickness mean : 26.9

```
[19]: # replace zero bloodpressure value with it's mean
print("Before BloodPressure mean : ",round(x_train.loc[:, 'BloodPressure'].
    ↪mean(),1))
x_test.loc[:, 'BloodPressure'] = x_test.loc[:, 'BloodPressure'].replace(0,
    ↪x_train.loc[:, 'BloodPressure'].mean())
x_train.loc[:, 'BloodPressure'] = x_train.loc[:, 'BloodPressure'].replace(0,
    ↪x_train.loc[:, 'BloodPressure'].mean())
print("After BloodPressure mean : ",round(x_train.loc[:, 'BloodPressure'].
    ↪mean(),1))
```

Before BloodPressure mean : 69.4

After BloodPressure mean : 72.8

```
[20]: # replace zero Glucose value with it's mean
print("Before Glucose mean : ",round(x_train.loc[:, 'Glucose'].mean(),1))
x_test.loc[:, 'Glucose'] = x_test.loc[:, 'Glucose'].replace(0, x_train.loc[:, 'Glucose'].mean())
x_train.loc[:, 'Glucose'] = x_train.loc[:, 'Glucose'].replace(0, x_train.loc[:, 'Glucose'].mean())
print("After Glucose mean : ",round(x_train.loc[:, 'Glucose'].mean(),1))
```

Before Glucose mean : 125.2
After Glucose mean : 126.1

```
[21]: # replace zero Insulin value with it's mean
print("Before Insulin mean : ",round(x_train.loc[:, 'Insulin'].mean(),1))
x_test.loc[:, 'Insulin'] = x_test.loc[:, 'Insulin'].replace(0, x_train.loc[:, 'Insulin'].mean())
x_train.loc[:, 'Insulin'] = x_train.loc[:, 'Insulin'].replace(0, x_train.loc[:, 'Insulin'].mean())
print("After Insulin mean : ",round(x_train.loc[:, 'Insulin'].mean(),1))
```

Before Insulin mean : 82.7
After Insulin mean : 123.8

3 Decision Tree

```
[22]: accuracy = {}
```

3.0.1 criterion="gini", splitter="best"

```
[23]: # Define and build model
clf = DecisionTreeClassifier(criterion="gini", splitter="best")
clf = clf.fit(x_train,y_train)
y_pred = clf.predict(x_test)
```

```
[24]: print(y_pred)
```

```
[1 1 1 0 0 1 0 1 0 0 0 0 1 1 1 1 0 0 0 1 0 1 0 1 1 0 1 0 0 1 0 1 0 0 0 0
 0 0 0 1 0 0 0 0 1 1 1 0 1 1 0 1 1 0 1 0 1 0 1 1 1 0 0 1 1 0 0 1 1 1
 1 1 0 0 0 0 0 1 0 1 1 1 0 0 1 0 0 1 0 0 1 1 0 0 0 0 0 1 0 1 0 0 0 1]
```

```
[25]: print(np.array(y_test))
```

```
[1 1 0 1 1 1 0 1 1 1 0 0 1 1 0 0 1 0 1 1 0 0 1 0 1 1 0 1 0 0 1 1 1 1 0 0 0
 1 0 0 1 1 0 0 0 1 0 0 0 1 1 1 0 0 0 1 1 1 0 1 0 1 1 0 0 0 1 1 0 0 1 1 1
 1 1 0 0 0 0 1 1 1 0 1 1 0 0 0 0 0 1 1 0 1 0 0 0 0 1 1 1 0 1 1 1 0 1]
```

```
[26]: accuracy["dt_gini_best"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.7129629629629629

```
[27]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[39 12]
 [19 38]]
```

```
[28]: print(metrics.classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.67	0.76	0.72	51
1	0.76	0.67	0.71	57
accuracy			0.71	108
macro avg	0.72	0.72	0.71	108
weighted avg	0.72	0.71	0.71	108

3.0.2 criterion="gini", splitter="best", max_depth=8

```
[29]: # Define and build model
      clf = DecisionTreeClassifier(criterion="gini", splitter="best", max_depth=8 )
      clf = clf.fit(x_train,y_train)
      y_pred = clf.predict(x_test)
```

```
[30]: print(y_pred)
```

```
[0 1 1 0 1 1 0 1 0 0 0 0 1 1 1 1 1 1 1 1 0 1 0 1 1 0 0 0 0 1 0 1 0 0 0 0
 1 0 0 1 0 0 1 0 1 1 1 0 1 0 0 1 0 0 1 1 0 1 0 1 0 1 1 0 0 0 0 1 0 0 1 1 1
 1 1 0 0 0 0 1 0 0 1 1 1 0 0 0 0 0 1 0 0 1 1 0 0 0 0 0 0 0 0 1 1 1 0 1]
```

```
[31]: print(np.array(y_test))
```

```
[1 1 0 1 1 1 0 1 1 1 0 0 1 1 0 0 1 0 1 1 0 0 1 0 0 1 1 1 1 0 0 0
 1 0 0 1 1 0 0 0 1 0 0 0 1 1 1 0 0 0 1 1 1 0 1 0 1 1 0 0 0 1 1 0 0 1 1 1
 1 1 0 0 0 0 1 1 1 0 1 1 0 0 0 0 0 1 1 0 1 0 0 0 0 1 1 1 0 1 1 1 0 1]
```

```
[32]: accuracy["dt_gini_best_8"] = metrics.accuracy_score(y_test, y_pred);
      print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.7314814814814815

```
[33]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[40 11]
 [18 39]]
```

```
[34]: print(metrics.classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.69	0.78	0.73	51

	1	0.78	0.68	0.73	57
accuracy				0.73	108
macro avg	0.73	0.73	0.73		108
weighted avg	0.74	0.73	0.73		108

3.0.3 criterion="entropy", splitter="best"

```
[35]: # Define and build model
      clf = DecisionTreeClassifier(criterion="entropy", splitter="best" )
      clf = clf.fit(x_train,y_train)
      y_pred = clf.predict(x_test)
```

```
[36]: print(y_pred)
```

```
[0 1 0 0 1 1 0 0 0 0 0 0 0 1 1 1 0 1 1 1 0 0 1 0 1 1 0 1 0 0 0 1 1 1 1 0 0
 1 0 1 0 0 1 0 0 1 1 1 0 1 1 0 0 1 0 0 0 1 1 0 1 1 1 1 0 0 1 1 1 0 0 1 1 0
 1 0 0 0 0 1 0 1 1 1 1 0 0 0 0 0 0 1 0 0 1 1 1 0 0 1 1 0 0 1 1 0 0 1]
```

```
[37]: print(np.array(y_test))
```

```
[1 1 0 1 1 1 0 1 1 1 0 0 1 1 0 0 1 0 1 1 0 0 1 0 1 1 0 1 0 0 1 1 1 1 0 0 0
 1 0 0 1 1 0 0 0 1 0 0 0 1 1 1 0 0 0 1 1 1 1 0 1 0 1 1 0 0 0 1 1 0 0 1 1 1
 1 1 0 0 0 0 1 1 1 0 1 1 0 0 0 0 0 1 1 0 1 0 0 0 0 1 1 1 0 1 1 1 0 1]
```

```
[38]: accuracy["dt_entropy_best"] = metrics.accuracy_score(y_test, y_pred);
      print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.6759259259259259

```
[39]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[36 15]
 [20 37]]
```

```
[40]: print(metrics.classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.64	0.71	0.67	51
1	0.71	0.65	0.68	57
accuracy			0.68	108
macro avg	0.68	0.68	0.68	108
weighted avg	0.68	0.68	0.68	108

3.0.4 criterion="entropy", splitter="best", max_depth=8

```
[41]: # Define and build model
      clf = DecisionTreeClassifier(criterion="entropy", splitter="best", max_depth=8 )
      clf = clf.fit(x_train,y_train)
      y_pred = clf.predict(x_test)
```

```
[42]: print(y_pred)
```

```
[0 1 0 0 1 1 0 0 0 0 0 0 1 1 1 1 0 1 1 1 0 0 0 0 1 1 0 1 0 1 1 1 1 0 1 0 0
 1 0 0 0 0 1 1 1 1 1 1 0 1 1 0 0 1 0 0 1 1 1 0 1 0 1 1 0 0 1 0 1 0 0 1 1 1
 1 0 0 0 1 1 0 1 1 1 1 1 0 0 1 0 0 1 0 0 1 1 1 0 0 1 1 0 0 1 1 1 0 1]
```

```
[43]: print(np.array(y_test))
```

```
[1 1 0 1 1 1 0 1 1 1 0 0 1 1 0 0 1 0 1 1 0 0 1 0 1 1 0 1 0 0 1 1 1 1 0 0 0
 1 0 0 1 1 0 0 0 1 0 0 0 1 1 1 0 0 0 1 1 1 1 0 1 0 1 1 0 0 0 1 1 0 0 1 1 1
 1 1 0 0 0 0 1 1 1 0 1 1 0 0 0 0 0 1 1 0 1 0 0 0 0 1 1 1 0 1 1 1 0 1]
```

```
[44]: accuracy["dt_entropy_best_8"] = metrics.accuracy_score(y_test, y_pred);
      print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.6759259259259259

```
[45]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[33 18]
 [17 40]]
```

```
[46]: print(metrics.classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.66	0.65	0.65	51
1	0.69	0.70	0.70	57
accuracy			0.68	108
macro avg	0.67	0.67	0.67	108
weighted avg	0.68	0.68	0.68	108

3.0.5 criterion="entropy", splitter="random"

```
[47]: # Define and build model
      clf = DecisionTreeClassifier(criterion="entropy", splitter="random" )
      clf = clf.fit(x_train,y_train)
      y_pred = clf.predict(x_test)
```

```
[48]: print(y_pred)
```

```
[0 1 1 0 1 1 0 1 1 1 0 1 1 1 1 1 1 1 1 1 0 0 1 1 0 0 0 0 0 0 1 1 0 0 1
 1 0 0 1 1 1 1 0 1 0 1 0 1 1 1 0 0 1 0 1 1 1 0 1 0 1 1 0 0 0 1 0 1 0 1 1 1
 1 0 0 0 0 0 1 0 1 0 0 1 0 0 1 0 0 1 1 0 0 0 0 0 0 0 0 0 1 0 1 0 0 0 1]
```

```
[49]: print(np.array(y_test))
```

```
[1 1 0 1 1 1 0 1 1 1 0 0 1 1 0 0 1 0 1 1 0 0 1 0 1 1 0 1 0 0 1 1 1 1 0 0 0
 1 0 0 1 1 0 0 0 1 0 0 0 1 1 1 0 0 0 1 1 1 1 0 1 0 1 1 1 0 0 0 1 1 0 0 1 1 1
 1 1 0 0 0 0 1 1 1 0 1 1 0 0 0 0 0 1 1 0 1 0 0 0 0 1 1 1 0 1 1 1 0 1]
```

```
[50]: accuracy["dt_entropy_random"] = metrics.accuracy_score(y_test, y_pred);
      print("Accuracy:", metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.7222222222222222

```
[51]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[37 14]
 [16 41]]
```

```
[52]: print(metrics.classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.70	0.73	0.71	51
1	0.75	0.72	0.73	57
accuracy			0.72	108
macro avg	0.72	0.72	0.72	108
weighted avg	0.72	0.72	0.72	108

3.0.6 criterion="entropy", splitter="random", max_depth=8

```
[53]: # Define and build model
      clf = DecisionTreeClassifier(criterion="entropy", splitter="random",
      ↪max_depth=8)
      clf = clf.fit(x_train, y_train)
      y_pred = clf.predict(x_test)
```

```
[54]: print(y_pred)
```

```
[1 1 1 0 1 1 0 1 1 1 1 0 1 1 0 1 1 1 1 1 0 0 1 0 0 0 0 0 0 0 1 1 1 1 1 0 0
 1 0 0 1 1 1 0 0 1 1 1 0 1 1 0 0 1 1 0 1 1 0 0 1 0 1 1 0 0 0 1 1 0 0 1 1 1
 1 0 0 0 1 0 1 0 0 1 1 1 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 1 0 1 1 0 0 1]
```

```
[55]: print(np.array(y_test))
```

```
[1 1 0 1 1 1 0 1 1 1 0 0 1 1 0 0 1 0 1 1 0 0 1 0 1 1 0 1 0 0 1 1 1 1 0 0 0
 1 0 0 1 1 0 0 0 1 0 0 0 1 1 1 0 0 0 1 1 1 1 0 1 0 1 1 1 0 0 0 1 1 0 0 1 1 1
 1 1 0 0 0 0 1 1 1 0 1 1 0 0 0 0 0 1 1 0 1 0 0 0 0 1 1 1 0 1 1 1 0 1]
```

```
[56]: accuracy["dt_entropy_random_8"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:", metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.75

```
[57]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[39 12]
 [15 42]]
```

```
[58]: print(metrics.classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.72	0.76	0.74	51
1	0.78	0.74	0.76	57
accuracy			0.75	108
macro avg	0.75	0.75	0.75	108
weighted avg	0.75	0.75	0.75	108

3.0.7 criterion="entropy", splitter="best", max_depth=3

```
[59]: # Define and build model
clf = DecisionTreeClassifier(criterion="entropy", splitter="best", max_depth=3 )
clf = clf.fit(x_train, y_train)
y_pred = clf.predict(x_test)
```

```
[60]: print(y_pred)
```

```
[1 1 0 0 1 1 1 1 0 1 0 0 0 1 1 1 1 1 1 0 0 1 0 1 1 1 1 0 1 1 1 1 1 0 0
 1 0 1 1 1 1 1 1 1 1 1 0 0 1 0 0 0 1 0 1 1 1 1 1 1 0 1 1 0 0 1 0 1 0 0 1 1 1
 1 1 0 0 1 1 1 1 1 1 1 1 0 0 1 0 0 0 1 0 1 1 0 0 0 0 0 1 0 1 1 1 0 1]
```

```
[61]: print(np.array(y_test))
```

```
[1 1 0 1 1 1 0 1 1 1 0 0 1 1 0 0 1 0 1 1 0 0 1 0 1 1 0 1 0 0 1 1 1 1 0 0 0
 1 0 0 1 1 0 0 0 1 0 0 0 1 1 1 0 0 0 1 1 1 1 0 1 0 1 1 0 0 0 1 1 0 0 1 1 1
 1 1 0 0 0 0 1 1 1 0 1 1 0 0 0 0 0 1 1 0 1 0 0 0 0 1 1 1 0 1 1 1 0 1]
```

```
[62]: accuracy["dt_entropy_best_3"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:", metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.7314814814814815

```
[63]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[31 20]
 [ 9 48]]
```


```
[64]: print(metrics.classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.78	0.61	0.68	51
1	0.71	0.84	0.77	57
accuracy			0.73	108
macro avg	0.74	0.72	0.72	108
weighted avg	0.74	0.73	0.73	108

```
[65]: feature_imp = pd.Series(clf.feature_importances_,index=X.columns).
      ↪sort_values(ascending=False)
print(feature_imp)
# Creating a bar plot
sns.barplot(x=feature_imp, y=feature_imp.index)
# Add labels to your graph
plt.xlabel('Feature Importance Score')
plt.ylabel('Features')
plt.title("Visualizing Important Features")
plt.legend()
plt.show()
```

Glucose	0.542718
BMI	0.256605
Age	0.155010
DiabetesPedigreeFunction	0.045667
Pregnancies	0.000000
BloodPressure	0.000000
SkinThickness	0.000000
Insulin	0.000000
dtype: float64	

No handles with labels found to put in legend.



3.0.8 criterion="entropy", splitter="random", max_depth=3

```
[66]: # Define and build model
      clf = DecisionTreeClassifier(criterion="entropy", splitter="random",
      ↪max_depth=3 )
      clf = clf.fit(x_train,y_train)
      y_pred = clf.predict(x_test)

[67]: print(y_pred)

[1 1 1 0 1 1 0 1 1 1 0 0 0 1 1 1 1 0 1 0 0 0 1 0 1 1 0 0 0 0 1 1 1 1 1 0 0
 1 0 0 1 1 1 0 0 1 1 1 0 1 1 0 1 1 0 1 0 1 0 0 1 0 1 1 0 0 1 1 1 0 0 1 1 1
 1 0 0 0 1 0 1 0 1 0 1 0 0 1 0 0 1 0 1 0 1 1 1 1 0 0 1 1 0 1 1 0 0 1]

[68]: print(np.array(y_test))

[1 1 0 1 1 1 0 1 1 1 0 0 1 1 0 0 1 0 1 1 0 0 1 0 1 1 0 1 0 0 1 1 1 1 0 0 0
 1 0 0 1 1 0 0 0 1 0 0 0 1 1 1 0 0 0 1 1 1 1 0 1 0 1 1 0 0 0 1 1 0 0 1 1 1
 1 1 0 0 0 0 1 1 1 0 1 1 0 0 0 0 0 1 1 0 1 0 0 0 0 1 1 1 0 1 1 1 0 1]

[69]: accuracy["dt_entropy_random_3"] = metrics.accuracy_score(y_test, y_pred);
      print("Accuracy:",metrics.accuracy_score(y_test, y_pred))

Accuracy: 0.7314814814814815

[70]: print(metrics.confusion_matrix(y_test, y_pred))

[[35 16]
 [13 44]]
```

```
[71]: print(metrics.classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.73	0.69	0.71	51
1	0.73	0.77	0.75	57
accuracy			0.73	108
macro avg	0.73	0.73	0.73	108
weighted avg	0.73	0.73	0.73	108

4 Accuracy visulization of Decision Tree

```
[72]: accuracy_df_dt = pd.DataFrame(list(zip(accuracy.keys(), accuracy.values())) ,
    ↪ columns = ['Arguments', 'Accuracy'])
accuracy_df_dt
```

```
[72]:
```

	Arguments	Accuracy
0	dt_gini_best	0.712963
1	dt_gini_best_8	0.731481
2	dt_entropy_best	0.675926
3	dt_entropy_best_8	0.675926
4	dt_entropy_random	0.722222
5	dt_entropy_random_8	0.750000
6	dt_entropy_best_3	0.731481
7	dt_entropy_random_3	0.731481

```
[73]: fig = px.bar(accuracy_df_dt, x='Arguments', y='Accuracy')
fig.show()
```

5 Random Forest

```
[74]: accuracy_rf = {}
```

5.0.1 n_estimators = 1000, criterion='entropy'

```
[75]: # Instantiate model with 1000 decision trees
rf = RandomForestClassifier(n_estimators = 1000, criterion='entropy' )
# Train the model on training data
rf.fit(x_train,y_train)
# Use the forest's predict method on the test data
y_pred = rf.predict(x_test)
```

```
[76]: print(y_pred)
```

```
[1 1 0 0 1 1 0 1 0 1 0 0 1 1 1 0 1 1 1 0 0 0 1 0 1 1 0 1 0 0 1 1 1 0 0 0 0
 1 0 0 1 1 1 1 0 1 0 0 0 1 1 0 0 1 0 1 1 1 0 0 1 0 1 1 0 0 1 0 1 0 0 1 1 0
 1 0 0 0 0 0 1 1 1 1 1 1 0 0 0 0 0 0 1 0 1 1 0 0 0 0 1 1 0 1 1 1 0 1]
```

```
[77]: print(np.array(y_test))
```

```
[1 1 0 1 1 1 0 1 1 1 0 0 1 1 0 0 1 0 1 1 0 0 1 0 1 1 0 1 0 0 1 1 1 1 0 0 0
 1 0 0 1 1 0 0 0 1 0 0 0 1 1 1 0 0 0 1 1 1 1 0 1 0 1 1 0 0 0 1 1 0 0 1 1 1
 1 1 0 0 0 0 1 1 1 0 1 1 0 0 0 0 0 1 1 0 1 0 0 0 0 1 1 1 0 1 1 1 0 1]
```

```
[78]: accuracy_rf["rf_entropy_1000"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:", metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.8240740740740741

```
[79]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[43  8]
 [11 46]]
```

```
[80]: print(metrics.classification_report(y_test, y_pred))
```

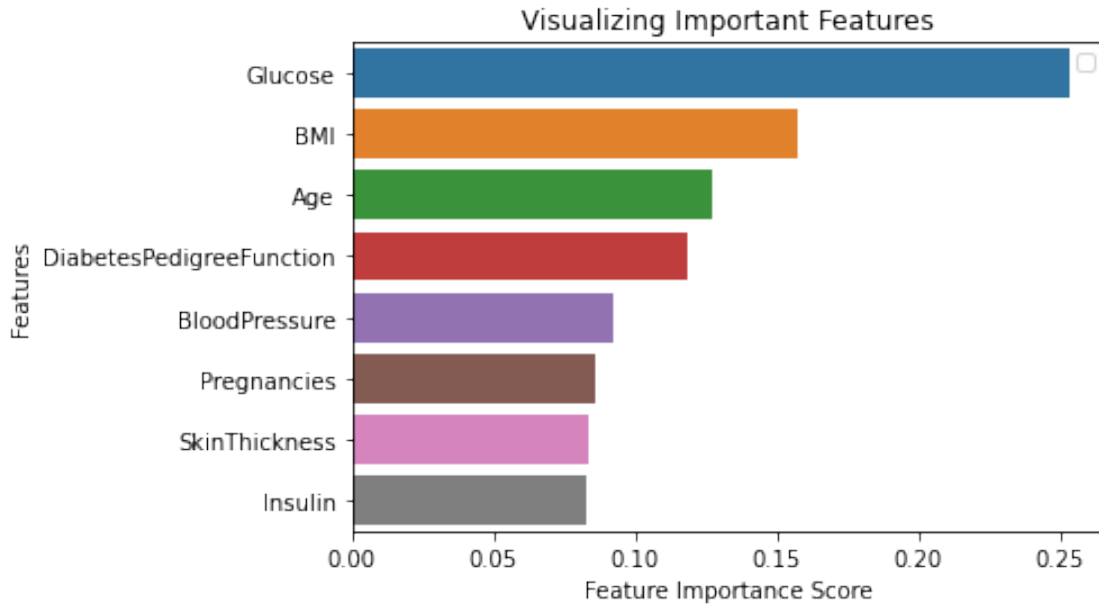
	precision	recall	f1-score	support
0	0.80	0.84	0.82	51
1	0.85	0.81	0.83	57
accuracy			0.82	108
macro avg	0.82	0.83	0.82	108
weighted avg	0.83	0.82	0.82	108

```
[81]: feature_imp = pd.Series(rf.feature_importances_, index=X.columns).
      ↪ sort_values(ascending=False)
print(feature_imp)
# Creating a bar plot
sns.barplot(x=feature_imp, y=feature_imp.index)
# Add labels to your graph
plt.xlabel('Feature Importance Score')
plt.ylabel('Features')
plt.title("Visualizing Important Features")
plt.legend()
plt.show()
```

No handles with labels found to put in legend.

Glucose	0.253201
BMI	0.157342
Age	0.127294
DiabetesPedigreeFunction	0.118516
BloodPressure	0.092122


```
Pregnancies      0.085738
SkinThickness    0.083264
Insulin          0.082523
dtype: float64
```



5.0.2 n_estimators = 100, criterion='entropy'

```
[82]: # Instantiate model with 100 decision trees
rf = RandomForestClassifier(n_estimators = 100, criterion='entropy' )
# Train the model on training data
rf.fit(x_train,y_train)
# Use the forest's predict method on the test data
y_pred = rf.predict(x_test)
```

```
[83]: print(y_pred)

[1 1 0 0 1 1 0 1 0 1 0 0 1 1 1 0 1 1 1 0 0 0 1 0 1 1 0 1 0 0 1 1 1 1 0 0 0
 1 0 0 1 1 1 1 1 0 1 0 0 0 1 1 0 0 1 0 1 1 1 0 0 1 0 1 1 0 0 1 0 1 0 0 1 1 0
 1 1 0 0 0 0 1 1 1 1 1 1 0 0 0 0 0 0 1 0 1 1 0 0 0 0 1 1 0 1 1 1 0 1]
```

```
[84]: print(np.array(y_test))

[1 1 0 1 1 1 0 1 1 1 0 0 1 1 0 0 1 0 1 1 0 0 1 0 1 1 0 1 0 0 1 1 1 1 0 0 0
 1 0 0 1 1 0 0 0 1 0 0 0 1 1 1 0 0 0 1 1 1 1 0 1 0 1 1 0 0 0 1 1 0 0 1 1 1
 1 1 0 0 0 0 1 1 1 0 1 1 0 0 0 0 0 1 1 0 1 0 0 0 0 1 1 1 0 1 1 1 0 1]
```

```
[85]: accuracy_rf["rf_entropy_100"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.8425925925925926

```
[86]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[43  8]
 [ 9 48]]
```

```
[87]: print(metrics.classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.83	0.84	0.83	51
1	0.86	0.84	0.85	57
accuracy			0.84	108
macro avg	0.84	0.84	0.84	108
weighted avg	0.84	0.84	0.84	108

5.0.3 n_estimators = 1000, random_state = 42, criterion='entropy'

```
[88]: # Instantiate model with 1000 decision trees
rf = RandomForestClassifier(n_estimators = 1000, random_state = 42,
    ↪criterion='entropy' )
# Train the model on training data
rf.fit(x_train,y_train)
# Use the forest's predict method on the test data
y_pred = rf.predict(x_test)
```

```
[89]: print(y_pred)
```

```
[1 1 0 0 1 1 0 1 0 1 0 0 1 1 0 0 1 1 1 0 0 0 1 0 1 1 0 1 0 0 1 1 1 1 0 0 0
 1 0 0 1 1 1 1 0 1 0 0 0 1 1 0 0 1 0 1 1 1 0 0 1 0 1 1 0 0 1 0 1 0 0 1 1 0
 1 0 0 0 0 0 1 1 1 1 1 1 0 0 0 0 0 0 1 0 1 1 0 0 0 0 1 1 0 1 1 1 0 1]
```

```
[90]: print(np.array(y_test))
```

```
[1 1 0 1 1 1 0 1 1 1 0 0 1 1 0 0 1 0 1 1 0 0 1 0 1 1 0 1 0 0 1 1 1 1 0 0 0
 1 0 0 1 1 0 0 0 1 0 0 0 1 1 1 0 0 0 1 1 1 1 0 1 0 1 1 0 0 0 1 1 0 0 1 1 1
 1 1 0 0 0 0 1 1 1 0 1 1 0 0 0 0 0 1 1 0 1 0 0 0 0 1 1 1 0 1 1 1 0 1]
```

```
[91]: accuracy_rf["rf_entropy_1000_42"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.8425925925925926

```
[92]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[44  7]
 [10 47]]
```

```
[93]: print(metrics.classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.81	0.86	0.84	51
1	0.87	0.82	0.85	57
accuracy			0.84	108
macro avg	0.84	0.84	0.84	108
weighted avg	0.84	0.84	0.84	108

5.0.4 n_estimators = 100, random_state = 42, criterion='entropy'

```
[94]: # Instantiate model with 100 decision trees
rf = RandomForestClassifier(n_estimators = 100, random_state = 42, max_depth = 8, criterion='entropy' )
# Train the model on training data
rf.fit(x_train,y_train)
# Use the forest's predict method on the test data
y_pred = rf.predict(x_test)
```

```
[95]: print(y_pred)
```

```
[1 1 0 0 1 1 0 1 0 1 0 0 1 1 0 0 1 1 1 0 0 0 1 0 1 1 0 1 0 0 1 1 1 0 0 0 0
 1 0 0 1 1 1 1 0 1 0 0 0 1 1 0 0 1 0 1 1 1 0 0 1 0 1 1 0 0 1 0 1 0 0 1 1 0
 1 0 0 0 0 0 1 1 1 1 1 1 0 0 0 0 0 1 1 0 1 1 0 0 0 0 1 1 0 1 1 1 0 1]
```

```
[96]: print(np.array(y_test))
```

```
[1 1 0 1 1 1 0 1 1 1 0 0 1 1 0 0 1 0 1 1 0 0 1 0 1 1 0 1 0 0 1 1 1 1 0 0 0
 1 0 0 1 1 0 0 0 1 0 0 0 1 1 1 0 0 0 1 1 1 1 0 1 0 1 1 0 0 0 1 1 0 0 1 1 1
 1 1 0 0 0 0 1 1 1 0 1 1 0 0 0 0 0 1 1 0 1 0 0 0 0 1 1 1 0 1 1 1 0 1]
```

```
[97]: accuracy_rf["rf_entropy_100_42"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.8425925925925926

```
[98]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[44  7]
 [10 47]]
```

```
[99]: print(metrics.classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.81	0.86	0.84	51
1	0.87	0.82	0.85	57

accuracy			0.84	108
macro avg	0.84	0.84	0.84	108
weighted avg	0.84	0.84	0.84	108

5.0.5 n_estimators = 1000, random_state = 42, max_depth = 8, criterion='entropy'

```
[100]: # Instantiate model with 1000 decision trees
rf = RandomForestClassifier(n_estimators = 1000, random_state = 42, max_depth = 8, criterion='entropy')
# Train the model on training data
rf.fit(x_train,y_train)
# Use the forest's predict method on the test data
y_pred = rf.predict(x_test)
```

```
[101]: print(y_pred)

[1 1 0 0 1 1 0 1 0 1 0 0 1 1 1 0 1 1 1 0 0 0 1 0 1 1 0 1 0 0 1 1 1 0 0 0 0
 1 0 0 1 1 1 1 0 1 0 0 0 1 1 0 0 1 0 1 1 1 0 0 1 0 1 1 0 0 1 0 1 0 0 1 1 0
 1 0 0 0 0 0 1 1 1 1 1 1 0 0 0 0 0 0 1 0 1 1 0 0 0 0 1 1 0 1 1 1 0 1]
```

```
[102]: print(np.array(y_test))

[1 1 0 1 1 1 0 1 1 1 0 0 1 1 0 0 1 0 1 1 0 0 1 0 1 1 0 1 0 0 1 1 1 1 0 0 0
 1 0 0 1 1 0 0 0 1 0 0 0 1 1 1 0 0 0 1 1 1 1 0 1 0 1 1 0 0 0 1 1 0 0 1 1 1
 1 1 0 0 0 0 1 1 1 0 1 1 0 0 0 0 0 1 1 0 1 0 0 0 0 1 1 1 0 1 1 1 0 1]
```

```
[103]: accuracy_rf["rf_entropy_1000_42_8"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.8240740740740741

```
[104]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[43  8]
 [11 46]]
```

```
[105]: print(metrics.classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.80	0.84	0.82	51
1	0.85	0.81	0.83	57
accuracy			0.82	108
macro avg	0.82	0.83	0.82	108
weighted avg	0.83	0.82	0.82	108

5.0.6 n_estimators = 100, random_state = 42, max_depth = 8, criterion='entropy'

```
[106]: # Instantiate model with 100 decision trees
rf = RandomForestClassifier(n_estimators = 100, random_state = 42, max_depth = 8, criterion='entropy' )
# Train the model on training data
rf.fit(x_train,y_train)
# Use the forest's predict method on the test data
y_pred = rf.predict(x_test)
```

```
[107]: print(y_pred)
```

```
[1 1 0 0 1 1 0 1 0 1 0 0 1 1 0 0 1 1 1 0 0 0 1 0 1 1 0 1 0 0 1 1 1 0 0 0 0
 1 0 0 1 1 1 1 0 1 0 0 0 1 1 0 0 1 0 1 1 1 0 0 1 0 1 1 0 0 1 0 1 0 0 1 1 0
 1 0 0 0 0 0 1 1 1 1 1 1 0 0 0 0 0 1 1 0 1 1 0 0 0 0 1 1 0 1 1 1 0 1]
```

```
[108]: print(np.array(y_test))
```

```
[1 1 0 1 1 1 0 1 1 1 0 0 1 1 0 0 1 0 1 1 0 0 1 0 1 1 0 1 0 0 1 1 1 1 0 0 0
 1 0 0 1 1 0 0 0 1 0 0 0 1 1 1 0 0 0 1 1 1 1 0 1 0 1 1 0 0 0 1 1 0 0 1 1 1
 1 1 0 0 0 0 1 1 1 0 1 1 0 0 0 0 0 1 1 0 1 0 0 0 0 1 1 1 0 1 1 1 0 1]
```

```
[109]: accuracy_rf["rf_entropy_100_42_8"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.8425925925925926

```
[110]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[44  7]
 [10 47]]
```

```
[111]: print(metrics.classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.81	0.86	0.84	51
1	0.87	0.82	0.85	57
accuracy			0.84	108
macro avg	0.84	0.84	0.84	108
weighted avg	0.84	0.84	0.84	108

5.0.7 n_estimators = 1000

```
[112]: # Instantiate model with 1000 decision trees
rf = RandomForestClassifier(n_estimators = 1000 )
# Train the model on training data
rf.fit(x_train,y_train)
```

```
# Use the forest's predict method on the test data
y_pred = rf.predict(x_test)
```

```
[113]: print(y_pred)
```

```
[1 1 0 0 1 1 0 1 0 1 0 0 1 1 1 0 1 1 1 0 0 0 1 0 1 1 0 1 0 0 1 1 1 0 0 0 0
 1 0 0 1 1 1 1 0 1 0 0 0 1 1 0 0 1 0 1 1 1 0 0 1 0 1 1 0 0 1 0 1 0 0 1 1 0
 1 0 0 0 0 0 1 1 1 1 1 1 0 0 0 0 0 0 1 0 1 1 0 0 0 0 1 1 0 1 1 1 0 1]
```

```
[114]: print(np.array(y_test))
```

```
[1 1 0 1 1 1 0 1 1 1 0 0 1 1 0 0 1 0 1 1 0 0 1 0 1 1 0 1 0 0 1 1 1 1 0 0 0
 1 0 0 1 1 0 0 0 1 0 0 0 1 1 1 0 0 0 1 1 1 1 0 1 0 1 1 0 0 0 1 1 0 0 1 1 1
 1 1 0 0 0 0 1 1 1 0 1 1 0 0 0 0 0 1 1 0 1 0 0 0 0 1 1 1 0 1 1 1 0 1]
```

```
[115]: accuracy_rf["rf_gini_1000"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:", metrics.accuracy_score(y_test, y_pred))
```

```
Accuracy: 0.8240740740740741
```

```
[116]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[43  8]
 [11 46]]
```

```
[117]: print(metrics.classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.80	0.84	0.82	51
1	0.85	0.81	0.83	57
accuracy			0.82	108
macro avg	0.82	0.83	0.82	108
weighted avg	0.83	0.82	0.82	108

5.0.8 n_estimators = 100

```
[118]: # Instantiate model with 100 decision trees
rf = RandomForestClassifier(n_estimators = 100 )
# Train the model on training data
rf.fit(x_train,y_train)
# Use the forest's predict method on the test data
y_pred = rf.predict(x_test)
```

```
[119]: print(y_pred)
```

```
[1 1 0 0 1 1 0 1 0 1 0 0 1 1 1 0 1 1 1 0 0 0 1 0 1 1 0 1 0 0 1 1 1 1 0 0 0
 1 0 0 1 1 1 1 0 1 0 0 0 1 1 0 0 1 0 1 1 1 0 0 1 0 1 1 0 0 1 0 1 0 0 1 1 0]
```

```
1 0 0 0 0 0 1 1 1 1 1 1 0 0 0 0 0 0 1 0 0 1 0 0 0 0 1 1 0 1 1 1 0 1]
```

```
[120]: print(np.array(y_test))
```

```
[1 1 0 1 1 1 0 1 1 1 0 0 1 1 0 0 1 0 1 1 0 0 1 0 1 1 0 1 0 0 1 1 1 1 0 0 0
 1 0 0 1 1 0 0 0 1 0 0 0 1 1 1 0 0 0 1 1 1 1 0 1 0 1 1 0 0 0 1 1 0 0 1 1 1
 1 1 0 0 0 0 1 1 1 0 1 1 0 0 0 0 0 1 1 0 1 0 0 0 0 1 1 1 0 1 1 1 0 1]
```

```
[121]: accuracy_rf["rf_gini_100"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:", metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.8240740740740741

```
[122]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[43  8]
 [11 46]]
```

```
[123]: print(metrics.classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.80	0.84	0.82	51
1	0.85	0.81	0.83	57
accuracy			0.82	108
macro avg	0.82	0.83	0.82	108
weighted avg	0.83	0.82	0.82	108

5.0.9 n_estimators = 1000, random_state = 42

```
[124]: # Instantiate model with 1000 decision trees
rf = RandomForestClassifier(n_estimators = 1000, random_state = 42 )
# Train the model on training data
rf.fit(x_train, y_train)
# Use the forest's predict method on the test data
y_pred = rf.predict(x_test)
```

```
[125]: print(y_pred)
```

```
[1 1 0 0 1 1 0 1 0 1 0 0 1 1 1 0 1 1 1 0 0 0 1 0 1 1 0 1 0 0 1 1 1 0 0 0 0
 1 0 0 1 1 1 1 0 1 0 0 0 1 1 0 0 1 0 1 1 1 0 0 1 0 1 1 0 0 1 0 1 0 0 1 1 0
 1 0 0 0 0 0 1 1 1 1 1 1 0 0 0 0 0 0 1 0 1 1 0 0 0 0 1 1 0 1 1 1 0 1]
```

```
[126]: print(np.array(y_test))
```

```
[1 1 0 1 1 1 0 1 1 1 0 0 1 1 0 0 1 0 1 1 0 0 1 0 1 1 0 1 0 0 1 1 1 1 0 0 0
 1 0 0 1 1 0 0 0 1 0 0 0 1 1 1 0 0 0 1 1 1 1 0 1 0 1 1 0 0 0 1 1 0 0 1 1 1
 1 1 0 0 0 0 1 1 1 0 1 1 0 0 0 0 0 1 1 0 1 0 0 0 0 1 1 1 0 1 1 1 0 1]
```

```
[127]: accuracy_rf["rf_gini_1000_42"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:", metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.8240740740740741

```
[128]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[43  8]
 [11 46]]
```

```
[129]: print(metrics.classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.80	0.84	0.82	51
1	0.85	0.81	0.83	57
accuracy			0.82	108
macro avg	0.82	0.83	0.82	108
weighted avg	0.83	0.82	0.82	108

5.0.10 n_estimators = 100, random_state = 42

```
[130]: # Instantiate model with 100 decision trees
rf = RandomForestClassifier(n_estimators = 100, random_state = 42, max_depth = 8 )
# Train the model on training data
rf.fit(x_train,y_train)
# Use the forest's predict method on the test data
y_pred = rf.predict(x_test)
```

```
[131]: print(y_pred)
```

```
[1 1 0 0 1 1 0 1 0 1 0 0 1 1 1 0 1 1 1 0 0 0 1 0 1 1 0 1 0 0 1 1 1 0 0 0 0
 1 0 0 1 1 1 1 0 1 0 0 0 1 1 0 0 1 0 1 1 1 0 0 1 0 1 1 0 0 1 0 1 0 0 1 1 0
 1 0 0 0 0 0 1 1 1 1 1 1 0 0 0 0 0 0 1 0 1 1 0 0 0 0 1 1 0 1 1 1 0 1]
```

```
[132]: print(np.array(y_test))
```

```
[1 1 0 1 1 1 0 1 1 1 0 0 1 1 0 0 1 0 1 1 0 0 1 0 1 1 0 1 0 0 1 1 1 1 0 0 0
 1 0 0 1 1 0 0 0 1 0 0 0 1 1 1 0 0 0 1 1 1 1 0 1 0 1 1 0 0 0 1 1 0 0 1 1 1
 1 1 0 0 0 0 1 1 1 0 1 1 0 0 0 0 0 1 1 0 1 0 0 0 0 1 1 1 0 1 1 1 0 1]
```

```
[133]: accuracy_rf["rf_gini_100_42"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:", metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.8240740740740741

```
[134]: print(metrics.confusion_matrix(y_test, y_pred))
```



```
[[43  8]
 [11 46]]
```

```
[135]: print(metrics.classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.80	0.84	0.82	51
1	0.85	0.81	0.83	57
accuracy			0.82	108
macro avg	0.82	0.83	0.82	108
weighted avg	0.83	0.82	0.82	108

5.0.11 n_estimators = 1000, random_state = 42, max_depth = 8

```
[136]: # Instantiate model with 1000 decision trees
rf = RandomForestClassifier(n_estimators = 1000, random_state = 42, max_depth = 8)
# Train the model on training data
rf.fit(x_train,y_train)
# Use the forest's predict method on the test data
y_pred = rf.predict(x_test)
```

```
[137]: print(y_pred)
```

```
[1 1 0 0 1 1 0 1 0 1 0 0 1 1 1 0 1 1 1 0 0 0 1 0 1 1 0 1 0 0 1 1 1 0 0 0 0
 1 0 0 1 1 1 1 0 1 0 0 0 1 1 0 0 1 0 1 1 1 0 0 1 0 1 1 0 0 1 0 1 0 0 1 1 0
 1 0 0 0 0 0 1 1 1 1 1 1 0 0 0 0 0 0 1 0 1 1 0 0 0 0 1 1 0 1 1 1 0 1]
```

```
[138]: print(np.array(y_test))
```

```
[1 1 0 1 1 1 0 1 1 1 0 0 1 1 0 0 1 0 1 1 0 0 1 0 1 1 0 1 0 0 1 1 1 1 0 0 0
 1 0 0 1 1 0 0 0 1 0 0 0 1 1 1 0 0 0 1 1 1 1 0 1 0 1 1 0 0 0 1 1 0 0 1 1 1
 1 1 0 0 0 0 1 1 1 0 1 1 0 0 0 0 0 1 1 0 1 0 0 0 0 1 1 1 0 1 1 1 0 1]
```

```
[139]: accuracy_rf["rf_gini_1000_42_8"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.8240740740740741

```
[140]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[43  8]
 [11 46]]
```

```
[141]: print(metrics.classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.80	0.84	0.82	51
1	0.85	0.81	0.83	57
accuracy			0.82	108
macro avg	0.82	0.83	0.82	108
weighted avg	0.83	0.82	0.82	108

5.0.12 n_estimators = 100, random_state = 42, max_depth = 8

```
[142]: # Instantiate model with 100 decision trees
rf = RandomForestClassifier(n_estimators = 100, random_state = 42, max_depth = 8)
# Train the model on training data
rf.fit(x_train,y_train)
# Use the forest's predict method on the test data
y_pred = rf.predict(x_test)
```

```
[143]: print(y_pred)

[1 1 0 0 1 1 0 1 0 1 0 0 1 1 1 0 0 0 1 0 1 1 0 1 0 0 1 1 1 0 0 0 0
 1 0 0 1 1 1 1 0 1 0 0 0 1 1 0 0 1 0 1 1 1 0 0 1 0 1 1 0 0 1 1 0
 1 0 0 0 0 0 1 1 1 1 1 1 0 0 0 0 0 0 1 0 1 1 0 0 0 0 1 1 0 1 1 1 0 1]
```

```
[144]: print(np.array(y_test))

[1 1 0 1 1 1 0 1 1 1 0 0 1 1 0 0 1 0 1 1 0 0 1 0 1 1 0 0 1 1 1 1 0 0 0
 1 0 0 1 1 0 0 0 1 0 0 0 1 1 1 0 0 0 1 1 1 1 0 1 0 1 1 0 0 0 1 1 0 0 1 1 1
 1 1 0 0 0 0 1 1 1 0 1 1 0 0 0 0 0 1 1 0 1 0 0 0 0 1 1 1 0 1 1 1 0 1]
```

```
[145]: accuracy_rf["rf_gini_100_42_8"] = metrics.accuracy_score(y_test, y_pred);
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.8240740740740741

```
[146]: print(metrics.confusion_matrix(y_test, y_pred))
```

```
[[43  8]
 [11 46]]
```

```
[147]: print(metrics.classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.80	0.84	0.82	51
1	0.85	0.81	0.83	57
accuracy			0.82	108
macro avg	0.82	0.83	0.82	108

weighted avg 0.83 0.82 0.82 108

6 Accuracy visulization of Random Forest

```
[148]: accuracy_df_rf = pd.DataFrame(list(zip(accuracy_rf.keys(), accuracy_rf.  
    ↪values()))), columns=['Arguments', 'Accuracy'])  
accuracy_df_rf
```

```
[148]:
```

	Arguments	Accuracy
0	rf_entropy_1000	0.824074
1	rf_entropy_100	0.842593
2	rf_entropy_1000_42	0.842593
3	rf_entropy_100_42	0.842593
4	rf_entropy_1000_42_8	0.824074
5	rf_entropy_100_42_8	0.842593
6	rf_gini_1000	0.824074
7	rf_gini_100	0.824074
8	rf_gini_1000_42	0.824074
9	rf_gini_100_42	0.824074
10	rf_gini_1000_42_8	0.824074
11	rf_gini_100_42_8	0.824074

```
[149]: fig = px.bar(accuracy_df_rf, x='Arguments', y='Accuracy')  
fig.show()
```

```
[150]: accuracy_df = pd.concat([accuracy_df_dt, accuracy_df_rf])  
accuracy_df['Accuracy'] = round(accuracy_df['Accuracy'] * 100, 2)  
fig = px.bar(accuracy_df, x='Arguments', y='Accuracy')  
print(accuracy_df['Accuracy'].max())  
fig.show()
```

84.26